

Lecture 19: Exploration and Exploitation

Bolei Zhou

The Chinese University of Hong Kong

bzhou@ie.cuhk.edu.hk

April 1, 2020

Outline

- 1 Introduction on exploration and exploitation
- 2 Multi-armed bandits
 - 1 Greedy and ϵ -greedy algorithms
 - 2 Temperature in Softmax bandit algorithm
 - 3 Upper Confidence Bound (UCB) algorithm
 - 4 Thompson sampling
- 3 Other exploration strategies in RL
 - 1 Entropy
 - 2 Curiosity
 - 3 Learning from failure

Trade-off between Exploration and Exploitation

- ① Decision making involves a fundamental choice:
Exploitation: Choose the best known action
Exploration: Explore some unknown action
- ② Short-term reward v.s. long-term reward: To collect information about action which results to long-term reward may involve the sacrifice in short-term reward.
- ③ Collect enough information to make the best overall decisions

Examples

① Going restaurant

- ① Exploitation: Go to your favorite restaurant
- ② Exploration: Try a new restaurant

② Oil Drilling

- ① Exploitation: Drill at the best known location
- ② Exploration: Drill at a new location

③ Webpage design

- ① Exploitation: Copy some existing template
- ② Exploration: Design your own from scratch

④ Game playing

- ① Exploitation: Play the move you already knows to work
- ② Exploration: Do some experimental move

⑤ New Year Resolution

- ① Exploitation: Stay in your comfort zone
- ② Exploration: Try some new thing

k -Armed Bandit



- 1 A multi-armed bandit is a tuple $\langle \mathcal{A}, \mathcal{R} \rangle$
- 2 k actions to take at each step t
- 3 $\mathcal{R}^a(r) = P(r|a)$ is unknown probability distribution over rewards
- 4 At each step t the agent selects an action $a_t \in \mathcal{A}$, then the environment generates a reward $r_t \sim \mathcal{R}^{a_t}$
- 5 The goal of agent is to maximize cumulative reward $\sum_{\tau=1}^T r_{\tau}$

Bernoulli Arm and Normal Arm

```
class BernoulliArm():
    def __init__(self, p):
        self.p = p
    def draw(self):
        if random.random() > self.p:
            return 0.0
        else:
            return 1.0

class NormalArm():
    def __init__(self, mu, sigma):
        self.mu = mu
        self.sigma = sigma
    def draw(self):
        return random.gauss(self.mu, self.sigma)
```

Definition of Value Function and Action-Value Function

- ① The action-value is the mean reward for action a (unknown)

$$Q(a) = \mathbb{E}(r|a) \quad (1)$$

- ② The optimal value

$$V^* = Q(a^*) = \max_{a \in \mathcal{A}} Q(a) \quad (2)$$

- ③ To estimate $Q(a)$, we can let $Q(a)$ at step t

$$Q_t(a) = \frac{\text{sum of rewards when } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t} = \frac{\sum_{i=1}^{t-1} R_i \cdot 1_{A_i=a}}{\sum_{i=1}^{t-1} 1_{A_i=a}} \quad (3)$$

Greedy Action and ϵ -Greedy Action to Take

- 1 The estimation of $Q(a)$ at step t

$$Q_t(a) = \frac{\text{sum of rewards when } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t} = \frac{\sum_{i=1}^{t-1} R_i \cdot 1_{A_i=a}}{\sum_{i=1}^{t-1} 1_{A_i=a}} \quad (4)$$

- 2 Greedy action selection algorithm: $A_t = \arg \max_a Q_t(a)$
- 3 Problem with the greedy algorithm?
- 4 ϵ -Greedy: greedy most of the time, but with small probability ϵ select random actions (ϵ is usually as 0.1)
 - 1 probability $1 - \epsilon$: $A_t = \arg \max_a Q_t(a)$
 - 2 probability ϵ : $A_t = \text{uniform}(\mathcal{A})$


ϵ -Greedy Algorithm

Algorithm 1 Simple *epsilon*-Greedy bandit algorithm

```
1: for  $a = 1$  to  $k$  do
2:    $Q(a) = 0, N(a) = 0$ 
3: end for
4: loop
5:    $A = \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \epsilon \\ \text{uniform}(\mathcal{A}) & \text{with probability } \epsilon \end{cases}$ 
6:    $R = \text{bandit}(A)$ 
7:    $N(A) = N(A) + 1$ 
8:    $Q(A) = Q(A) + \frac{1}{N(A)}[R - Q(A)]$ 
9: end loop
```

Optimistic Initial Values

- 1 Simple idea: initialize $Q(a)$ to high value
- 2 Encourage the exploration over all possible actions early on



Softmax Bandit Algorithm

- ① To learn a numerical preference for each action a (like learning a policy function, denoted as $H_t(a)$,

$$\pi_t(A_t) = P(A_t = a) = \frac{e^{H_t(a)}}{\sum_{b=1}^k e^{H_t(b)}} \quad (5)$$

- ② Learning based on the idea of stochastic gradient descend

$$\text{For } A_t, H_{t+1}(A_t) = H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)). \quad (6)$$

$$\text{For all } a \neq A_t, H_{t+1}(a) = H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a). \quad (7)$$

here \bar{R}_t is the average of all the rewards up to time t . Page 38 contains the full derivation.

- ③ Learning based on the incremental estimation

$$H_t(A_t) = H_{t-1}(A_t) + \frac{1}{N_t(A_t)}[R - H_{t-1}(A_t)] \quad (8)$$

Temperature

- ① Scaling factor, temperature τ , to control the degree of exploration. High temperature, atoms will behave more random.

$$P(A_t = a) = \frac{e^{H_t(a)/\tau}}{\sum_{b=1}^k e^{H_t(b)/\tau}} \quad (9)$$

Annealing

- ① Annealing is the process of modifying an algorithm's behavior so that it will explore less over time
- ② Effect to different algorithms:
 - ① To reduce ϵ in ϵ -Greedy algorithm
 - ② To make the temperature go lower and lower in Softmax Bandit algorithm

UCB - The Upper Confidence Bound Algorithm

- ① Both the ϵ -Greedy algorithm and the Softmax algorithm share the following broad properties:
 - ① select the arm that currently has the highest estimated value
 - ② explore and choose an arm that isn't the one that currently seems the best
 - ③ reduce the exploration by annealing (dynamically change some parameters ϵ and τ)
- ② UCB takes a very different approach.
 - ① UCB does not use randomness at all
 - ② UCB doesn't have any free parameters to configure before you can deploy it.

UCB - The Upper Confidence Bound Algorithm

- 1 $U_t(a)$ is the upper confidence bound of the reward value, so that the true value is below the bound with **high probability**,

$$Q(a) \leq Q_t(a) + U_t(a) \quad (10)$$

- 2 The upper bound $U_t(a)$ is a function of $N_t(a)$, where a larger number of trials $N_t(a)$ should give us a smaller bound $N_t(a)$ (less uncertain).
- 3 In UCB1 algorithm,


$$U_t(a) = \sqrt{\frac{2 \log t}{N_t(a)}} \quad (11)$$

- 4 Thus the action is selected as to maximize the UCB

$$a_t = \arg \max_a \left[Q_t(a) + \sqrt{\frac{2 \log t}{N_t(a)}} \right] \quad (12)$$

UCB - The Upper Confidence Bound Algorithm

- 1 If we are uncertain about an action, we should optimistically assume that it is the correct action.



UCB - The Upper Confidence Bound Algorithm

$$a_t = \arg \max_a [Q_t(a) + \sqrt{\frac{2 \log t}{N_t(a)}}] \quad (13)$$

- 1 Upper bound term brings significant values from the beginning.
- 2 UCB is an explicitly curiosity-driven algorithm that tries to seek out the unknown
- 3 Square root term is a measure of the uncertainty or variance in the estimate of a 's value. Each time a is selected, the uncertainty term decreases.
- 4 Logarithm increases get smaller over time.

Deriving the Upper Confidence Bound

Hoeffding's Inequality

Let X_1, \dots, X_n be i.i.d random variables and they are all bounded by the interval $[0, 1]$. The sample mean is $\bar{X}_n = \frac{1}{n} \sum_{\tau=1}^n X_\tau$. Then for $u > 0$, we have:

$$P(\mathbb{E}[X] > \bar{X}_n + u) \leq e^{-2nu^2} \quad (14)$$

- 1 Following the Hoeffding's Inequality, then we have

$$P(Q(a) > Q_t(a) + U_t(a)) \leq e^{-2N_t(a)U_t(a)^2}$$

Deriving the Upper Confidence Bound

- 1 We want to pick up a bound so that with high chances the true mean is below the sample mean + the upper confidence bound,


$$P(Q(a) > Q_t(a) + U_t(a)) \leq e^{-2N_t(a)U_t(a)^2} = p,$$

$$\text{thus, } U_t(a) = \sqrt{\frac{-\log p}{2N_t(a)}}$$

- 2 One heuristic is to reduce the threshold p in time, as we want to make more confident bound estimation with more rewards observed. Set $p = t^{-4}$ we get UCB1 algorithm:

$$a_t^{UCB1} = \arg \max_a [Q_t(a) + \sqrt{\frac{2 \log t}{N_t(a)}}] \quad (15)$$

Example Code for Bandit Algorithms



- Bandit Algorithms for Website Optimization by John Myles White.
- <https://github.com/cuhkrlcourse/RLexample/tree/master/bandits>

Thompson Sampling: Bayesian decision making

- 1 An algorithm for online decision problems where actions are taken sequentially
- 2 Bayesian inference to compute the posterior with the known prior and the likelihood of getting the sampled data

Algorithm 1 BernGreedy(K, α, β)

```
1: for  $t = 1, 2, \dots$  do
2:   #estimate model:
3:   for  $k = 1, \dots, K$  do
4:      $\hat{\theta}_k \leftarrow \alpha_k / (\alpha_k + \beta_k)$ 
5:   end for
6:
7:   #select and apply action:
8:    $x_t \leftarrow \operatorname{argmax}_k \hat{\theta}_k$ 
9:   Apply  $x_t$  and observe  $r_t$ 
10:
11:  #update distribution:
12:   $(\alpha_{x_t}, \beta_{x_t}) \leftarrow (\alpha_{x_t}, \beta_{x_t}) + (r_t, 1 - r_t)$ 
13: end for
```

Algorithm 2 BernThompson(K, α, β)

```
1: for  $t = 1, 2, \dots$  do
2:   #sample model:
3:   for  $k = 1, \dots, K$  do
4:     Sample  $\hat{\theta}_k \sim \text{beta}(\alpha_k, \beta_k)$ 
5:   end for
6:
7:   #select and apply action:
8:    $x_t \leftarrow \operatorname{argmax}_k \hat{\theta}_k$ 
9:   Apply  $x_t$  and observe  $r_t$ 
10:
11:  #update distribution:
12:   $(\alpha_{x_t}, \beta_{x_t}) \leftarrow (\alpha_{x_t}, \beta_{x_t}) + (r_t, 1 - r_t)$ 
13: end for
```

- 3 More detailed tutorial: <https://arxiv.org/pdf/1707.02038.pdf>

Other Exploration Strategies in RL

- ① Entropy-regularized policy optimization
- ② Learning from internal rewards: Curiosity-driven exploration
- ③ Learning from failures: Hindsight Experience Replay

Entropy-regularized policy optimization


- ① SAC incorporates **entropy regularization**
- ② Entropy is a quantity which measures how random a random variable is, $H(P) = E_{x \sim P}[-\log P(x)]$
- ③ Entropy-regularized RL: the policy is trained to maximize a trade-off between expected return and entropy, a measure of randomness in the policy

$$\pi^* = \arg \max E_{\tau \sim \pi} \left[\sum_t \gamma^t (R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot | s_t))) \right]$$

Curiosity-driven Learning


- 1 Environments with sparse rewards or non-existing rewards are very challenging

Vizdoom "DoomMyWayHome"



From Felix Steger

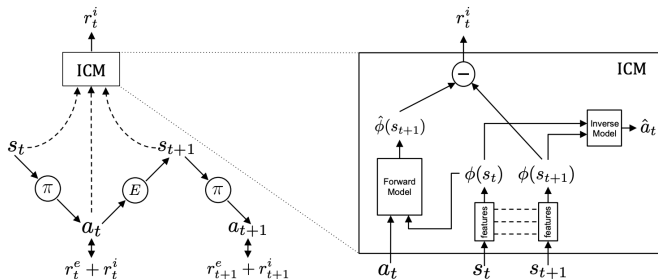
Montezuma's Revenge



- 2 A new reward function: curiosity
 - 1 Curiosity is an intrinsic reward which can be considered as the error of the agent to predict the consequence of its own actions given its current state
 - 2 Encourage the agent to perform actions that reduce the uncertainty in the agent's ability to predict the consequence of its own action
 - 3 How to measure the curiosity

Curiosity-driven Learning

- 1 Pathak ICML'17 Curiosity-driven Exploration by self-supervised prediction: Intrinsic Curiosity Module (ICM)




- 1 Forward model is to predict the transition dynamics (world model):
$$\hat{\phi}_{s_{t+1}} = f(\phi(s_t), a_t; \theta_f)$$
- 2 Inverse Model is trained to predict the action a given s_t and s_{t+1} :
$$\hat{a} = g(\phi(s_t), \phi(s_{t+1}); \theta_g)$$
- 3 Intrinsic reward is defined as residual $r_t^i = \|\hat{\phi}(s_{t+1}) - \phi(s_{t+1})\|$

Curiosity-driven Learning


- ① Pathak ICML'17 Curiosity-driven Exploration by self-supervised prediction: Intrinsic Curiosity Module (ICM)
- ② Demo: <https://pathak22.github.io/noreward-rl/>
- ③ Further work: Large-Scale Study of Curiosity-Driven Learning. ICLR'19:
<https://pathak22.github.io/large-scale-curiosity/>

Goal-oriented Environments

Pick and place at a desired goal





Manipulate block to a desired goal



- 1 The desired goal might appear at any state
- 2 The reward is very sparse


Goal-oriented Environments

① NIPS'17 Hindsight Experience Replay (HER): Learning from failure



- ① Intuitive idea: Even though we have not succeeded at a specific goal, we have at least achieved a different one.
- ② So we can just pretend that we wanted to achieve this goal to begin with, instead of the original one
- ③ By doing this substitution, the reinforcement learning algorithm can obtain a learning signal since it has achieved some goal (**so we create some dense reward**)

HER for Goal-oriented Environments



Summary for Exploration in RL

- ① It is beneficial to create extra **external rewards** and **internal rewards** for the agent to learn