

Lecture 22: Generative Modeling

Bolei Zhou

The Chinese University of Hong Kong

Yann LeCun's Cake

■ "Pure" Reinforcement Learning (cherry)

- ▶ The machine predicts a scalar reward given once in a while.
- ▶ **A few bits for some samples**

■ Supervised Learning (icing)

- ▶ The machine predicts a category or a few numbers for each input
- ▶ Predicting human-supplied data
- ▶ **10→10,000 bits per sample**

■ Unsupervised/Predictive Learning (cake)

- ▶ The machine predicts any part of its input for any observed part.
- ▶ Predicts future frames in videos
- ▶ **Millions of bits per sample**

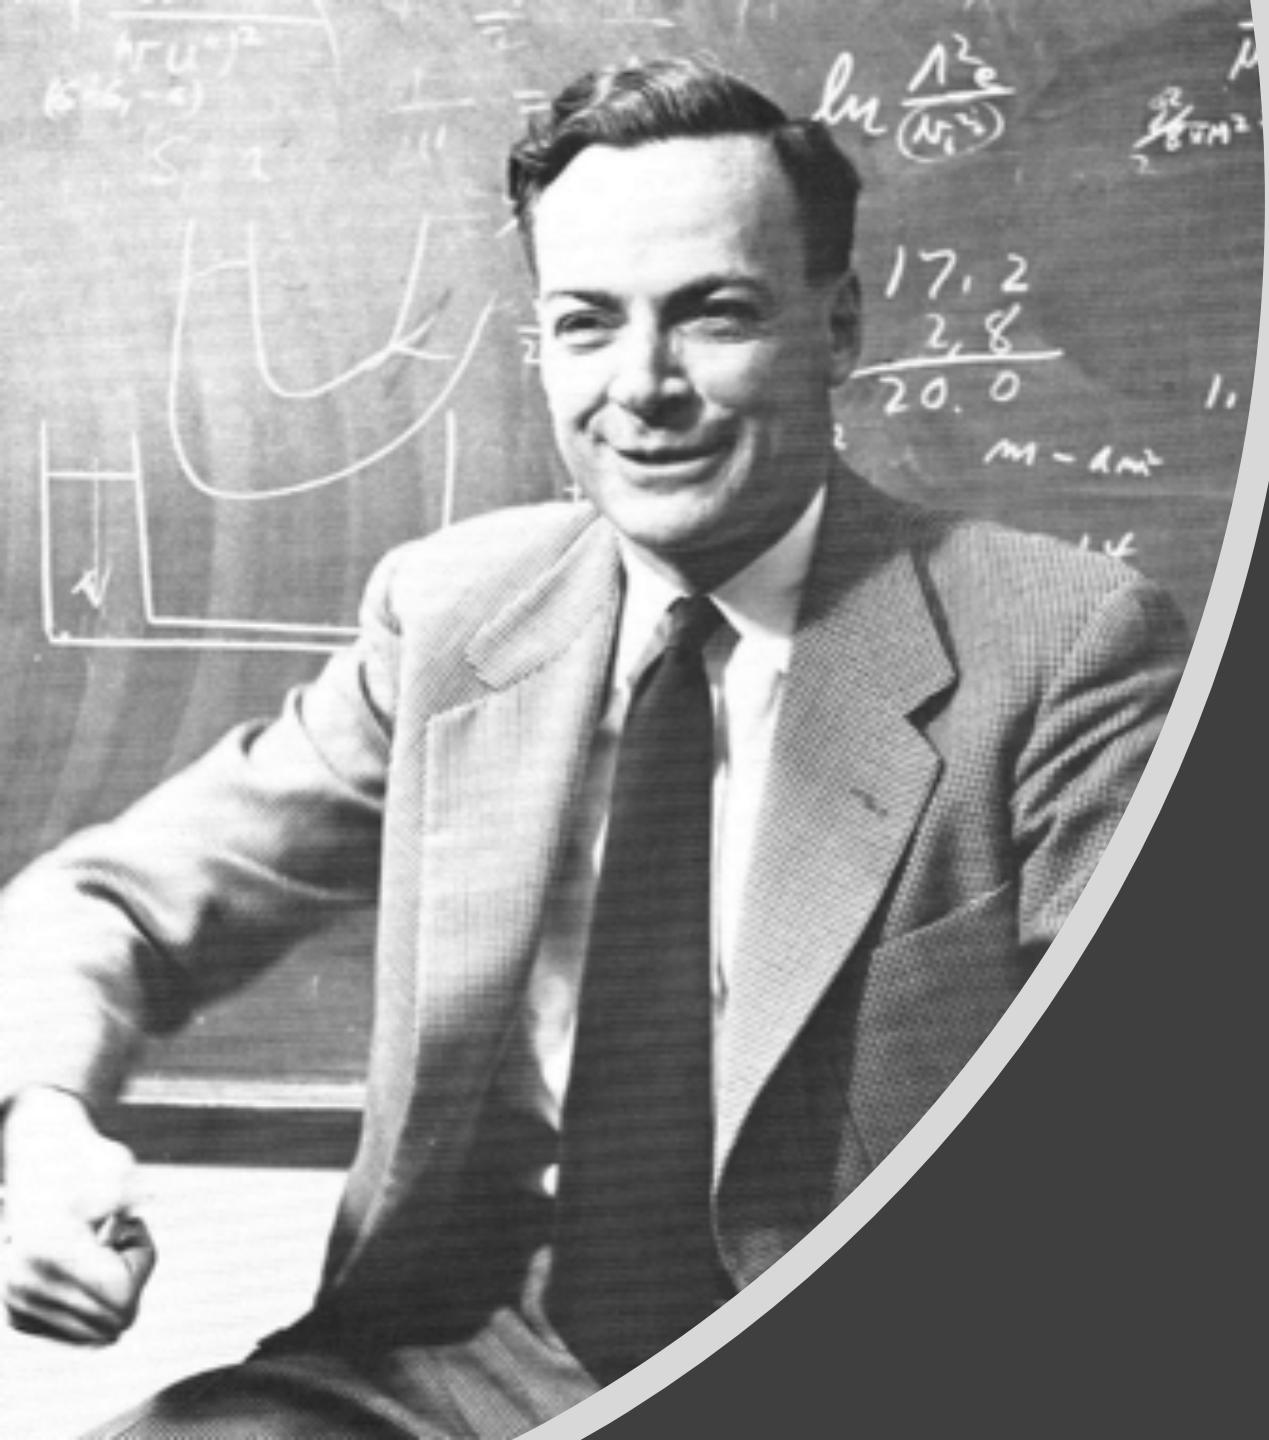


Let's remove the fancy
cherry and icings

Unsupervised Learning

Generative
modeling

Self-supervised
feature learning



Richard Feynman:
“What I cannot create,
I do not understand”

Inference and Generation



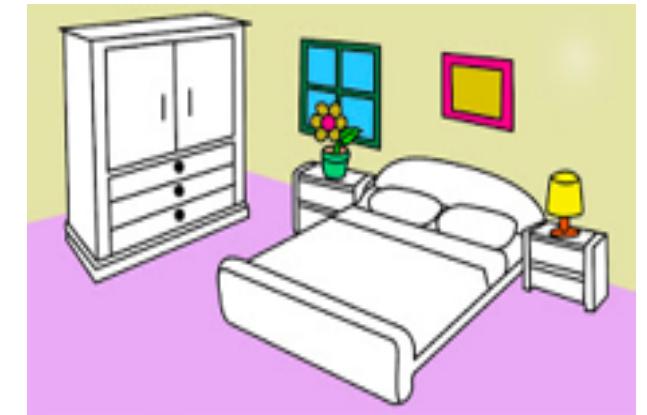
Inference →



Predictions:
bedroom
wood
foliage

layout

Generation →



Statistical Generative Models

- Data: samples such as images of bedrooms
- Prior knowledge: distribution functions (Gaussian, or other non-linear?), loss function (e.g., maximum likelihood, adversarial loss?),



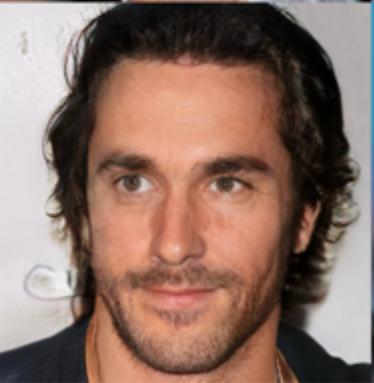
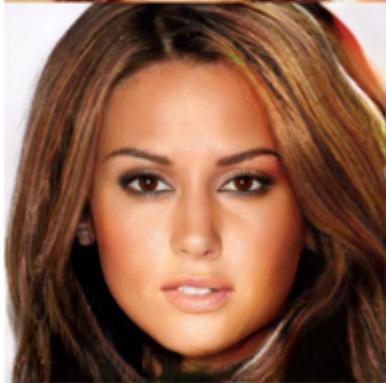
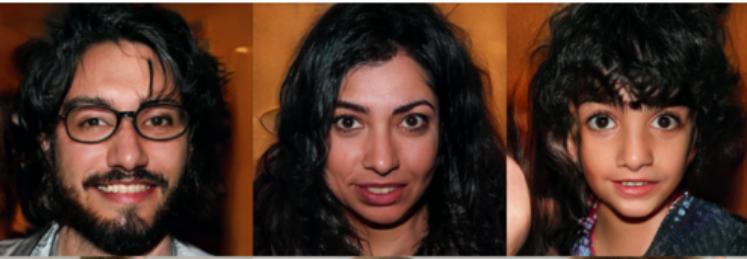
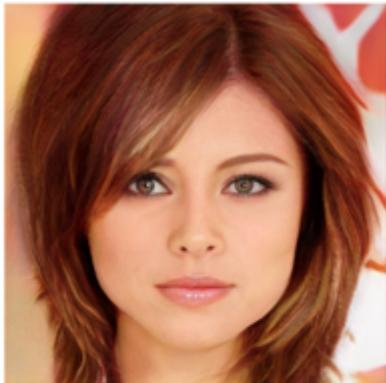
learning
→

learned probability
distribution $p(x)$

sampling
→



Progress for Image Generation



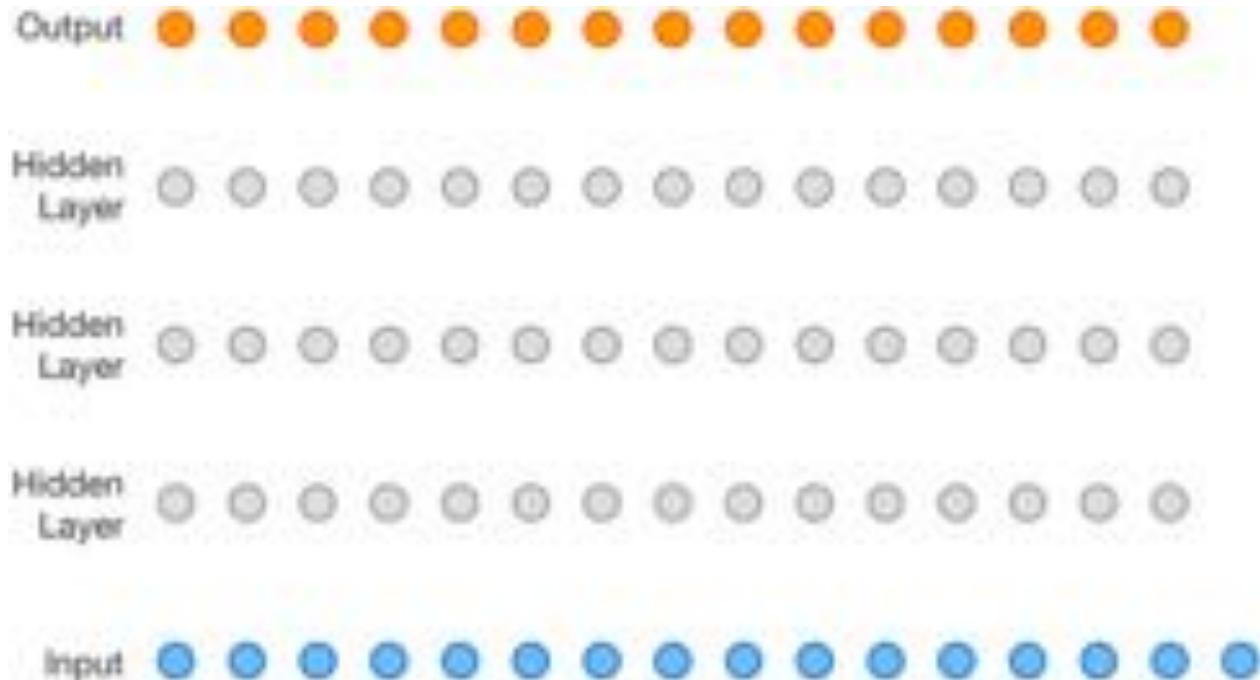
PG-GAN

StyleGAN

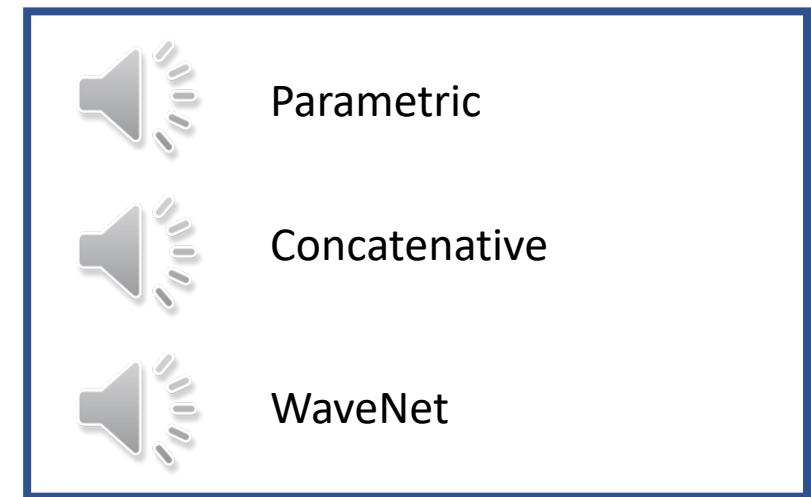
BigGAN

Progress for Speech Synthesis

WaveNet



Text to Speech



van den Oord et al, 2016c

Image Super Resolution

Conditional generative model $P(\text{high res image} \mid \text{low res image})$



Ledig et al., 2017

Machine Translation

Conditional generative model $P(\text{ English text} | \text{ Chinese text})$

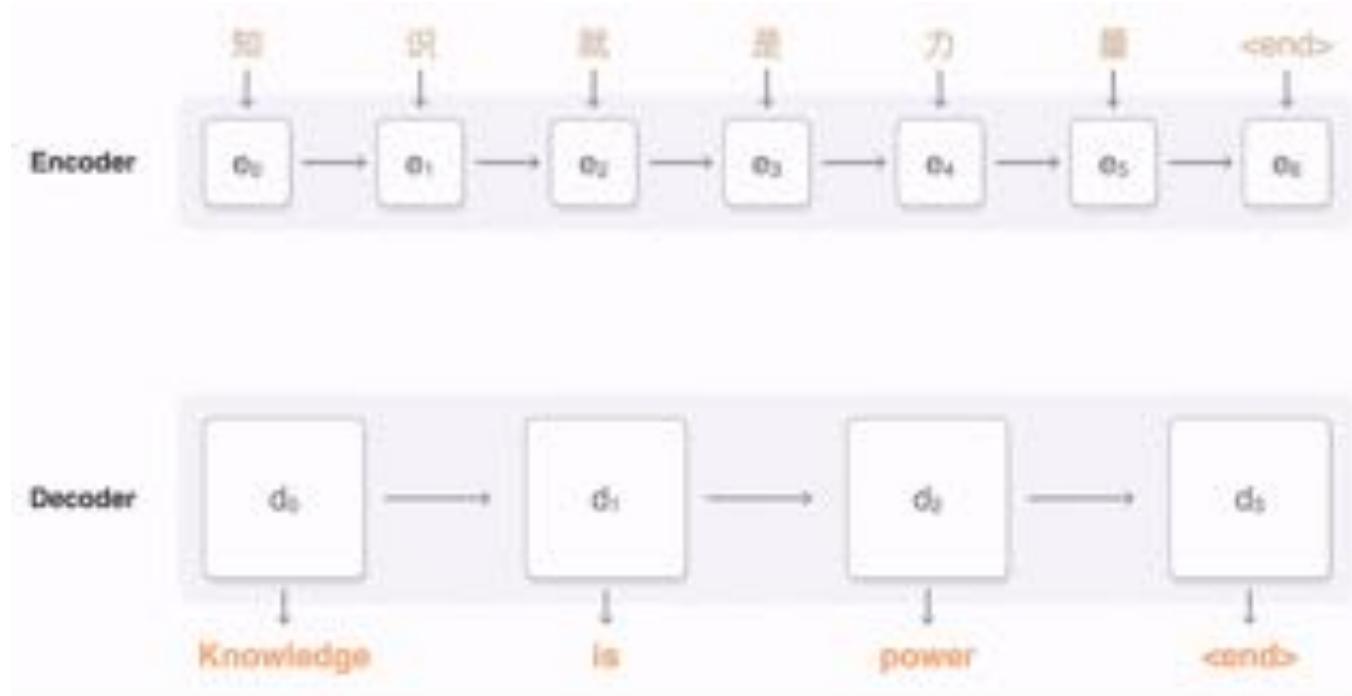


Figure from Google AI research blog.

Paragraph Generation

Custom prompt

To better learn statistics, students should



Completion

To better learn statistics, students should understand the functions of probability, statistics, and probability distributions. They should also understand the basic logic that underlies statistics and find their own ideas and methods for applying the logic to new problems. Their basic understanding of probability and statistics should provide the foundation they will need to think creatively and critically when dealing with critical situations and problems. They should understand the information we can find out by running statistical tests, and they should understand the implications and limitations of such information. They should be able to apply basic math concepts, statistics, probability and mathematics to their own particular circumstances and to teach other students. They should be able to analyze and critique the many issues of statistics that are in our society. They should be able to tell about possible sources

DeepFakes (video+audio)

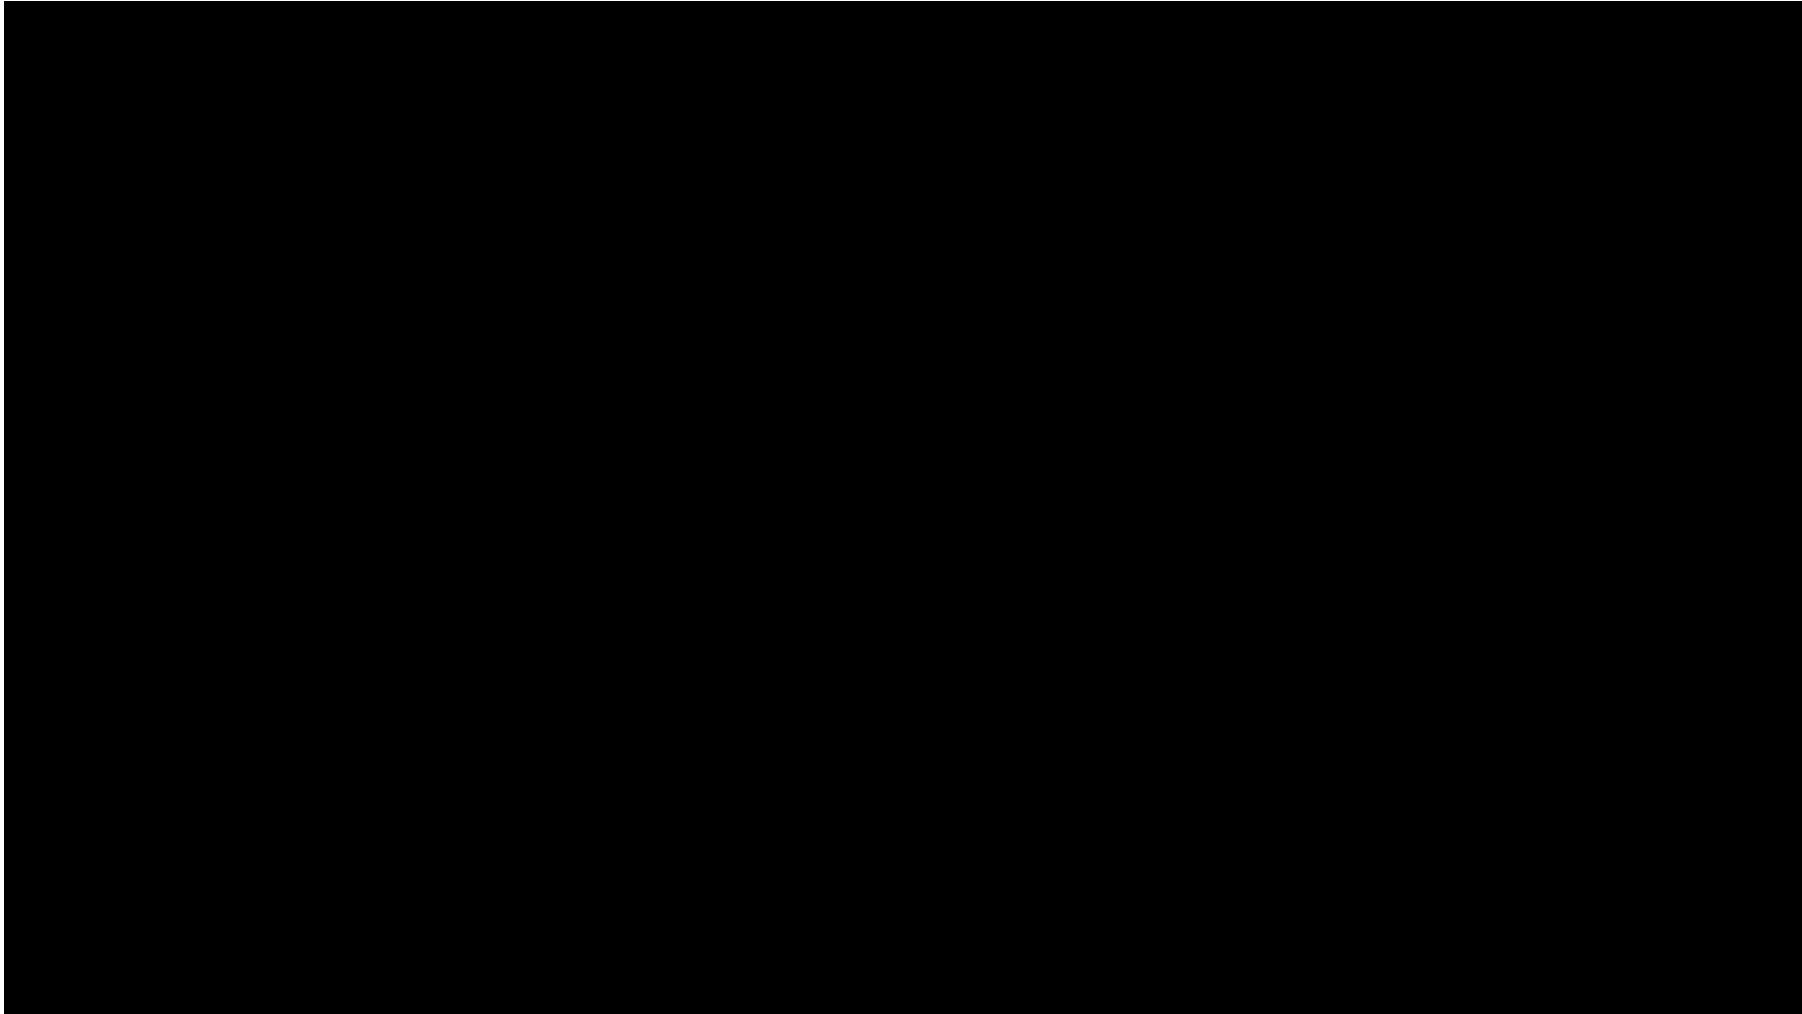


Image Translation

Conditional generative model $P(\text{ zebra images} | \text{ horse images})$



Zhu et al., 2017

(Very) Brief look on generative models

- Likelihood-based models: Autoregressive models
- Latent variable models: Variational Autoencoder
- Implicit generative models: Generative Adversarial Networks (GANs)
- Latent semantics in GANs

Learning a generative model

- We want to learn a probability distribution $p(x)$ over images x such that we can sample it to generate x



learning
→

learned probability
distribution $p(x)$

sampling
→



Autoregressive generative model

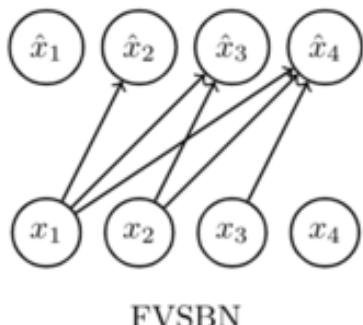
- Given a dataset of handwritten digits (binarized MNIST)



- Each image has $n = 28 \times 28 = 784$ pixels, with each pixel either being black (0) or white (1)
- Goal: Learn a probability distribution $p(x) = p(x_1, \dots, x_{784})$

Autoregressive generative model

- We can use chain rule for factorization:
 - $p(x_1, \dots, x_{784}) = p(x_1)p(x_2|x_1)p(x_3|x_1,x_2)\dots p(x_n|x_1,\dots,x_{n-1})$
- Then we can have the following simple autoregressive model
 - $P(x_1=1;a_1) = a_1, P(x_1=0;a_1) = 1-a_1$
 - $P(x_2=1|x_1;a_2) = \text{logit}(a_{02}+a_{12}x_1)$
 - $P(x_3=1|x_1,x_2;a_3) = \text{logit}(a_{03}+a_{13}x_1+a_{23}x_2)$

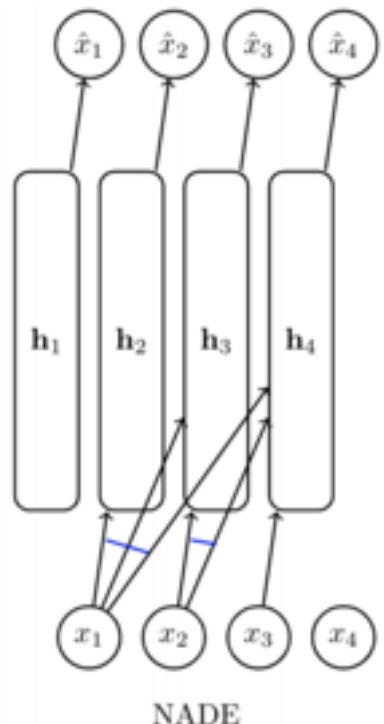


$$\hat{x}_i = p(X_i = 1|x_1, \dots, x_{i-1}; \boldsymbol{\alpha}^i) = p(X_i = 1|x_{<i}; \boldsymbol{\alpha}^i) = \sigma(\alpha_0^i + \sum_{j=1}^{i-1} \alpha_j^i x_j)$$

Fully visible sigmoid belief network (FVSBN)

NADE: Neural Autoregressive Density Estimation

- To improve the factorization model: use one-layer neural network instead of logistic regression



$$\hat{x}_i = p(x_i | x_1, \dots, x_{i-1}; \underbrace{A_i, \mathbf{c}_i, \boldsymbol{\alpha}_i, b_i}_{\text{parameters}}) = \sigma(\boldsymbol{\alpha}_i \mathbf{h}_i + b_i)$$
$$\mathbf{h}_i = \sigma(A_i \mathbf{x}_{<i} + \mathbf{c}_i)$$

$$\mathbf{h}_2 = \sigma \left(\underbrace{\begin{pmatrix} \vdots \\ \textcolor{blue}{w}_1 \\ \vdots \end{pmatrix}}_{A_2} x_1 \right) \quad \mathbf{h}_3 = \sigma \left(\underbrace{\begin{pmatrix} \vdots & \vdots \\ \textcolor{blue}{w}_1 & \textcolor{red}{w}_2 \\ \vdots & \vdots \end{pmatrix}}_{A_3} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \right) \quad \mathbf{h}_4 = \sigma \left(\underbrace{\begin{pmatrix} \vdots & \vdots & \vdots \\ \textcolor{blue}{w}_1 & \textcolor{red}{w}_2 & \textcolor{blue}{w}_3 \\ \vdots & \vdots & \vdots \end{pmatrix}}_{A_3} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} \right)$$

Training data



Samples from the probability



Figure from The Neural Autoregressive Distribution Estimator, 2011.

RNN: Recurrent Neural Networks

Suppose $x_i \in \{h, e, l, o\}$. Use one-hot encoding:

- h encoded as $[1, 0, 0, 0]$, e encoded as $[0, 1, 0, 0]$, etc.

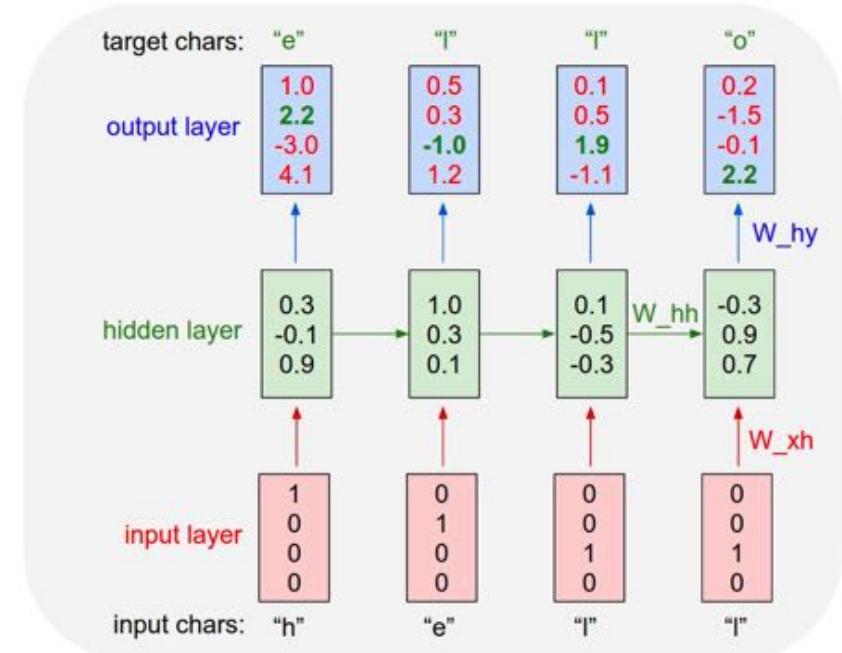
Autoregressive: $p(x = \text{hello}) = p(x_1 = h)p(x_2 = e|x_1 = h)p(x_3 = l|x_1 = h, x_2 = e) \cdots p(x_5 = o|x_1 = h, x_2 = e, x_3 = l, x_4 = l)$

For example,

$$p(x_2 = e|x_1 = h) = \text{softmax}(o_1) = \frac{\exp(2.2)}{\exp(1.0) + \cdots + \exp(4.1)}$$

$$o_1 = W_{hy} h_1$$

$$h_1 = \tanh(W_{hh} h_0 + W_{xh} x_1)$$



RNN: Recurrent Neural Networks

RNN model on Wikipedia

Naturalism and decision for the majority of Arab countries' capitalide was grounded by the Irish language by [[John Clair]], [[An Imperial Japanese Revolt]], associated with Guangzham's sovereignty. His generals were the powerful ruler of the Portugal in the [[Protestant Immineners]], which could be said to be directly in Cantonese Communication, which followed a ceremony and set inspired prison, training. The emperor travelled back to [[Antioch, Perth, October 25|21]] to note, the Kingdom of Costa Rica, unsuccessful fashioned the [[Thrales]], [[Cynth's Dajoard]], known in western [[Scotland]], near Italy to the conquest of India with the conflict. Copyright was the succession of independence in the slop of Syrian influence that was a famous German movement based on a more popular servicious, non-doctrinal and sexual power post. Many governments recognize the military housing of the [[Civil Liberalization and Infantry Resolution 265 National Party in Hungary]], that is sympathetic to be to the [[Punjab Resolution]]
(PJS)[<http://www.humah.yahoo.com/guardian>.

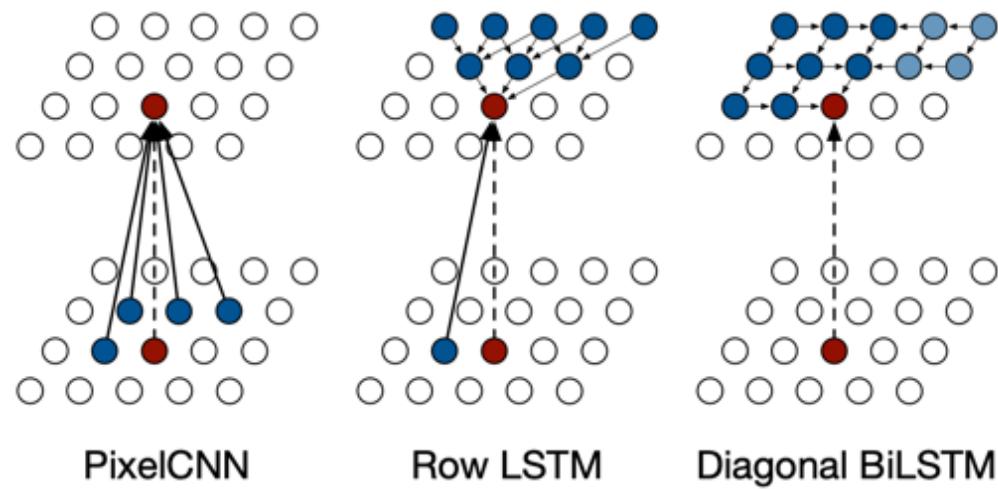
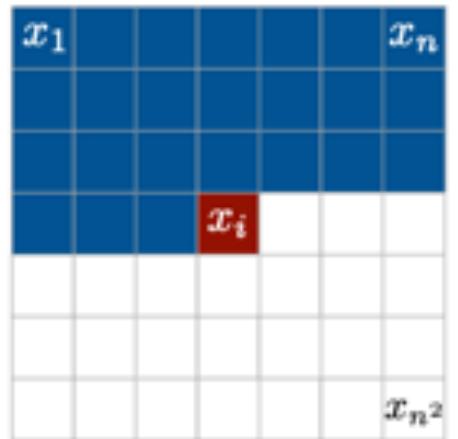
RNN model on Linux source code

```
/*
 * Increment the size file of the new incorrect UI_FILTER group information
 * of the size generatively.
 */
static int indicate_policy(void)
{
    int error;
    if (fd == MARN_EPT) {
        /*
         * The kernel blank will coeld it to userspace.
         */
        if (ss->segment < mem_total)
            unblock_graph_and_set_blocked();
        else
            ret = 1;
        goto bail;
    }
    segaddr = in_SB(in.addr);
    selector = seg / 16;
    setup_works = true;
    for (i = 0; i < blocks; i++) {
        seq = buf[i++];
        bpf = bd->bd.next + i * search;
        if (fd) {
            current = blocked;
        }
    }
}
```

PixelRNN and PixelCNN (Oord et al. 2016)

- Model image pixel by pixel using raster scan order

$$p(x_t \mid x_{1:t-1}) = p(x_t^{\text{red}} \mid x_{1:t-1})p(x_t^{\text{green}} \mid x_{1:t-1}, x_t^{\text{red}})p(x_t^{\text{blue}} \mid x_{1:t-1}, x_t^{\text{red}}, x_t^{\text{green}})$$

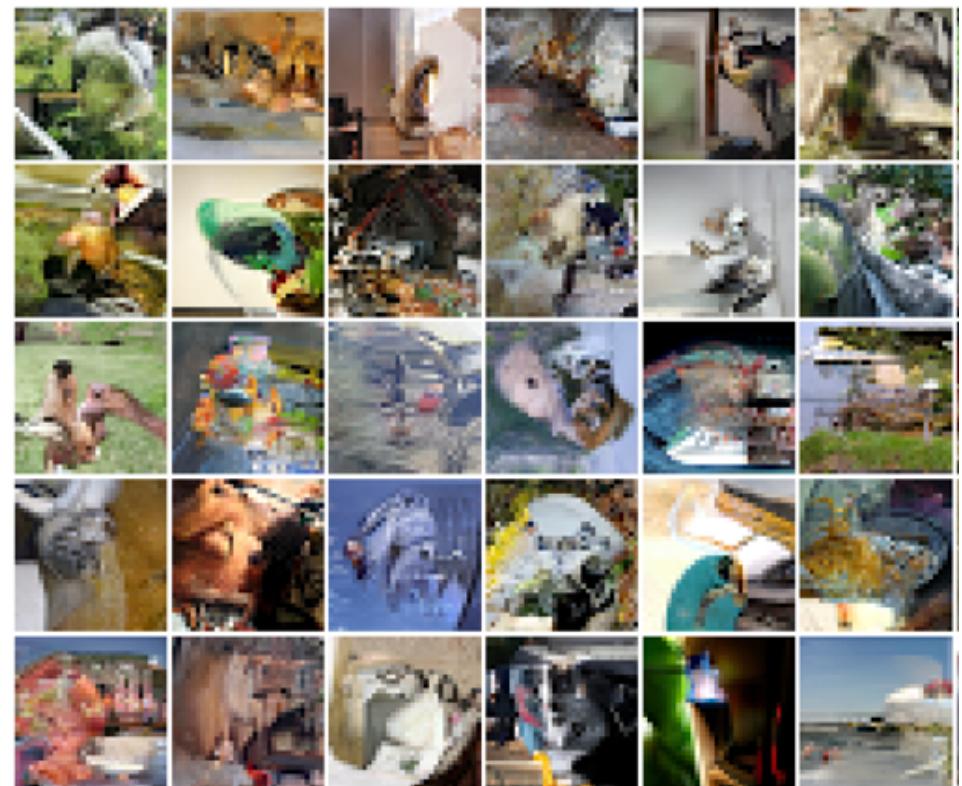


PixelCNN and PixelRNN (Oord et al. 2016)



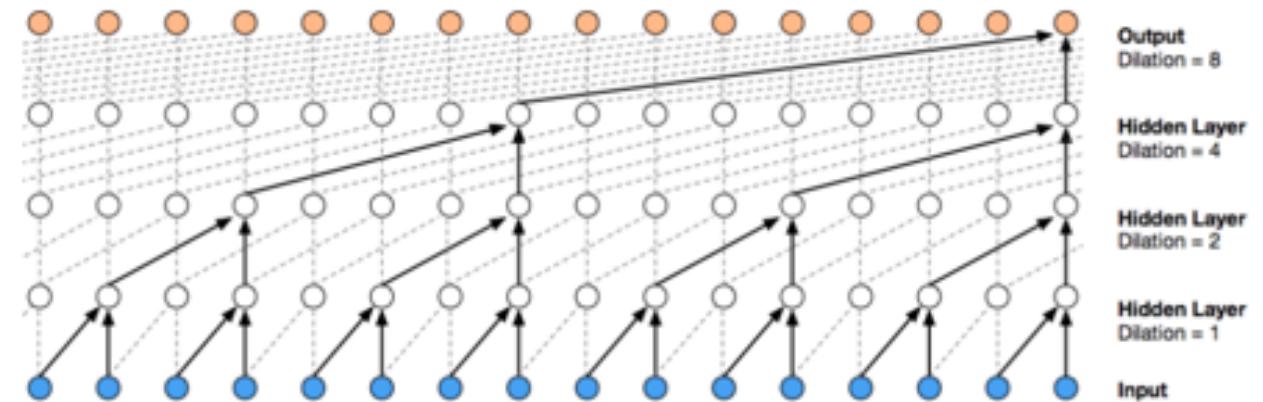
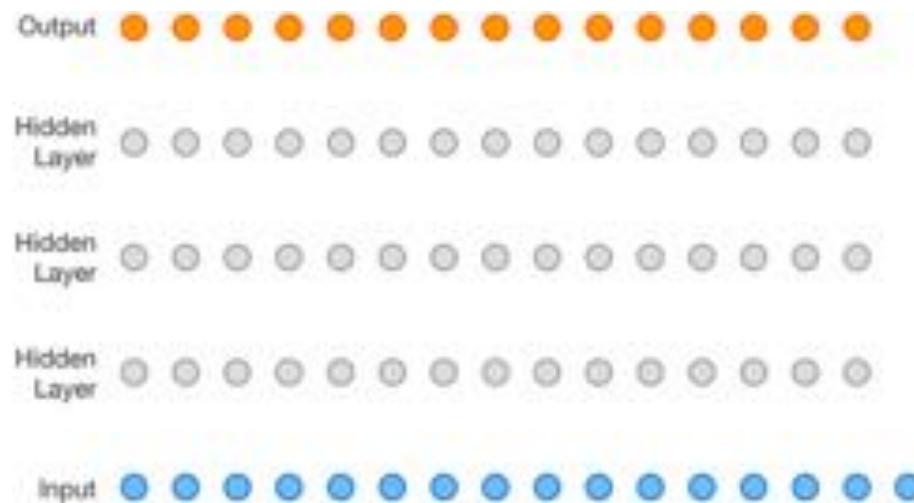
Figure 1. Image completions sampled from a PixelRNN.

Samples from PixelRNN trained on ImageNet



WaveNet (Oord et al. 2016)

- CNN with 1D dilated convolutions



Summary of Autoregressive Models

- Autoregressive models:
 - Chain-rule factorization is fully general
 - Compact representation via conditional independence and/or neural parameterizations
- Pros:
 - Easy to evaluate likelihoods
 - Easy to train
- Cons:
 - Requires an ordering
 - Generation is sequential
 - Cannot learn features in an unsupervised way

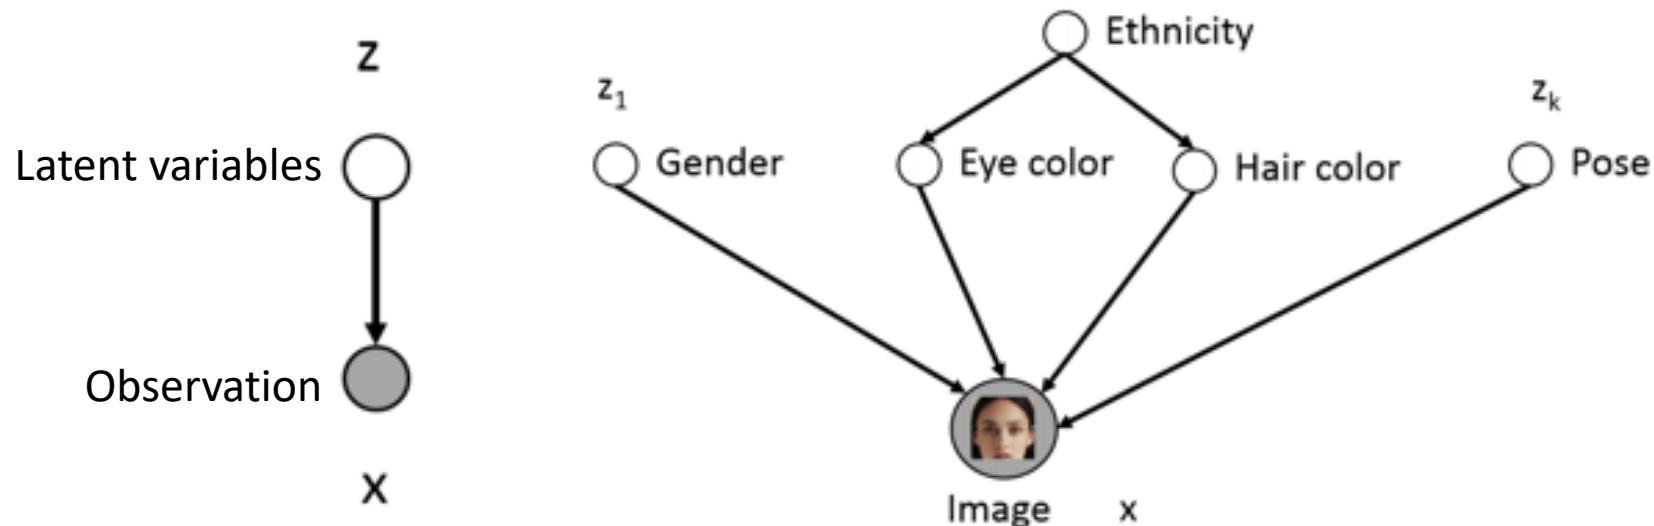
Latent Variable Models

- Gender, eye color, race, pose, etc are the variation factors in images, unless annotated, they are not explicitly available
- Then we can model them as (low-dimensional) **latent** variables



Latent Variable Models

Data generation process could be driven by the latent variables



Latent Variable Models

- We can define complex models $p(x)$ in terms of simple building blocks $p(x|z)$
- It can be applied for unsupervised learning tasks easily
- Challenge: much more difficult to learn due to the latent variables.

A mixture of an infinite number of Gaussians:

① $\mathbf{z} \sim \mathcal{N}(0, I)$

② $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mu_\theta(\mathbf{z}), \Sigma_\theta(\mathbf{z}))$ where $\mu_\theta, \Sigma_\theta$ are neural networks

③ Even though $p(\mathbf{x} | \mathbf{z})$ is simple, the marginal $p(\mathbf{x})$ is very complex/flexible

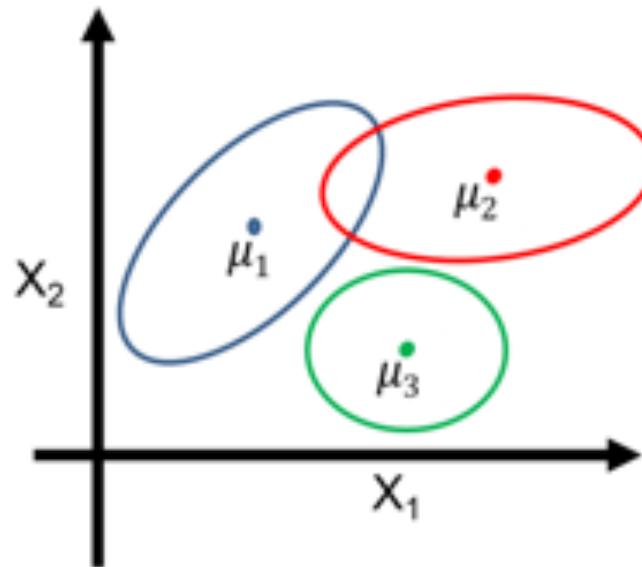
Mixture of Gaussians as a simple latent variable model

Mixture of Gaussians. Bayes net: $\mathbf{z} \rightarrow \mathbf{x}$.

- ① $\mathbf{z} \sim \text{Categorical}(1, \dots, K)$
- ② $p(\mathbf{x} | \mathbf{z} = k) = \mathcal{N}(\mu_k, \Sigma_k)$

Generative process:

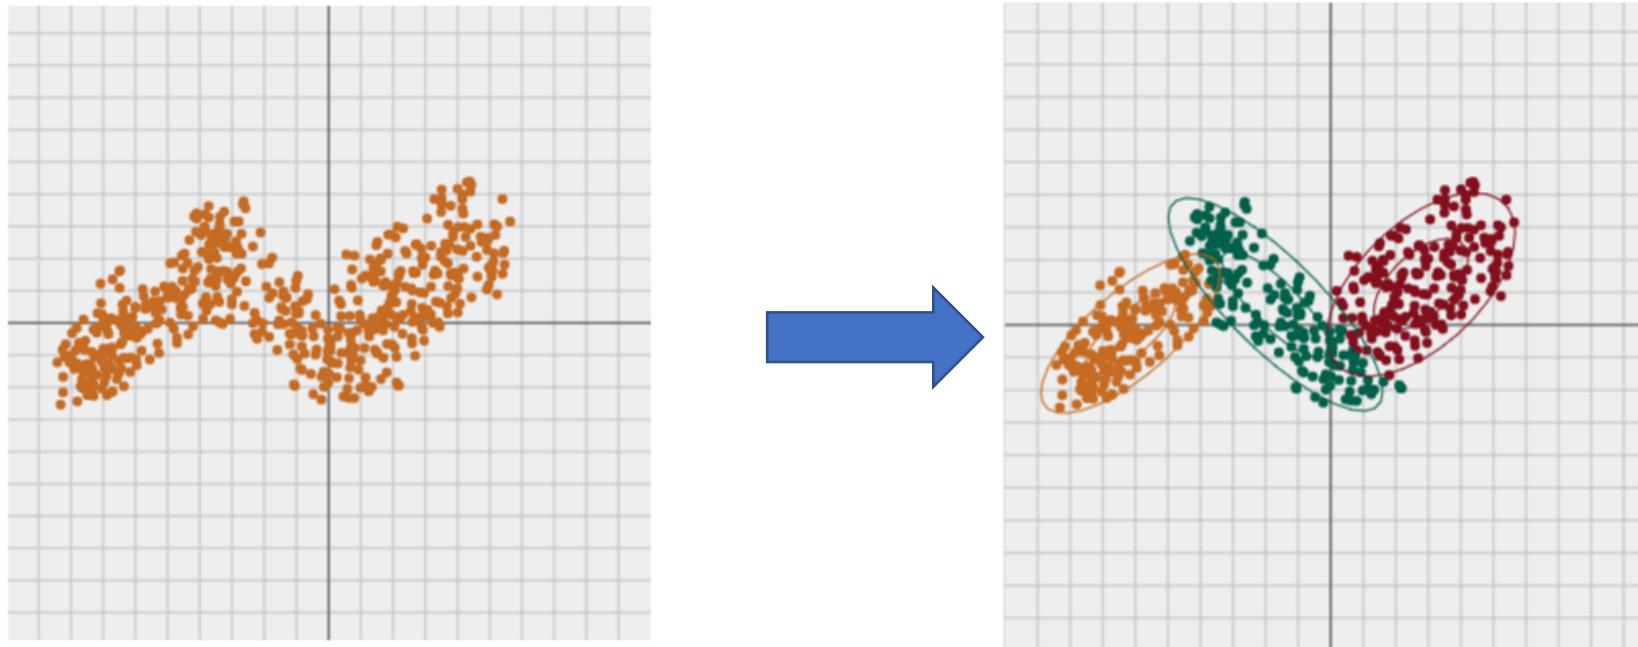
- Pick a mixture component k by sampling \mathbf{z}
- Generate a data point by sampling from that Gaussian



Mixture of Gaussians as a simple latent variable model

- Unsupervised learning: the posterior $p(z|x)$ identifies the mixture component

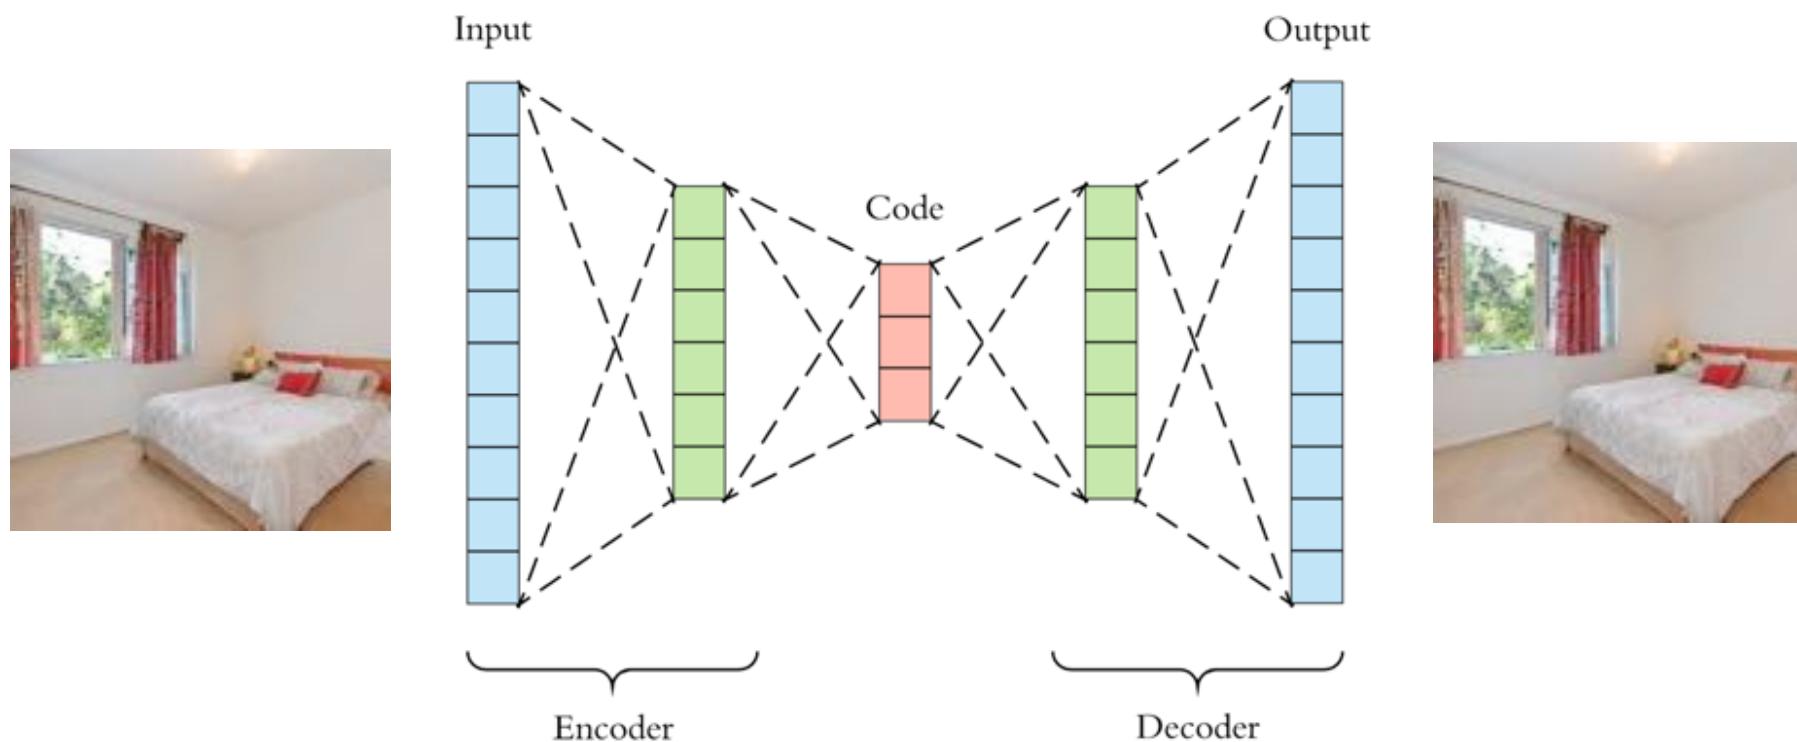
$$i = \operatorname{argmax} p(z=i|x)$$



$$p(\mathbf{x}) = \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z}) = \sum_{\mathbf{z}} p(\mathbf{z})p(\mathbf{x} | \mathbf{z}) = \sum_{k=1}^K p(\mathbf{z} = k) \underbrace{\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}_{\text{component}}$$

Autoencoder

- Autoencoder is originally proposed for dimensionality reduction and feature learning

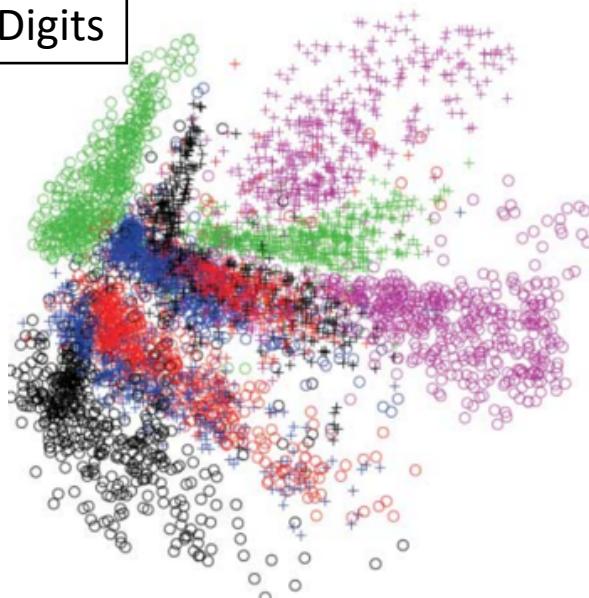


Autoencoder

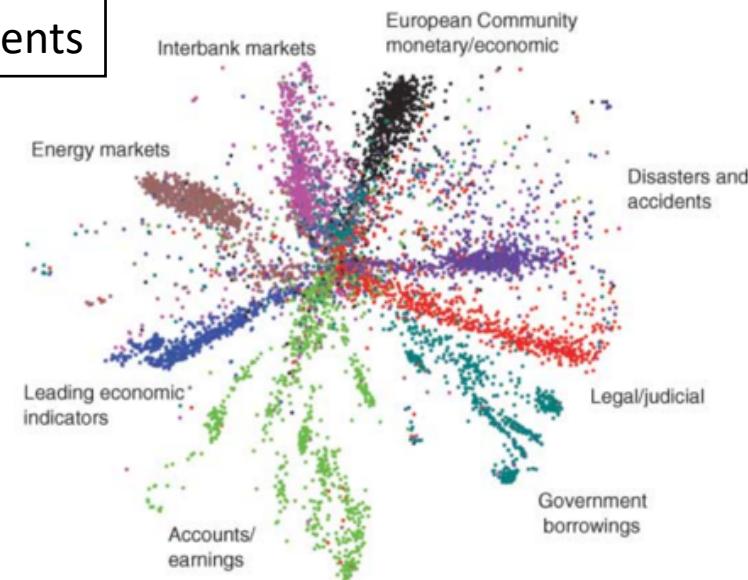
- Autoencoder is originally proposed for dimensionality reduction and feature learning

Clustering using the latent features

MNIST Digits

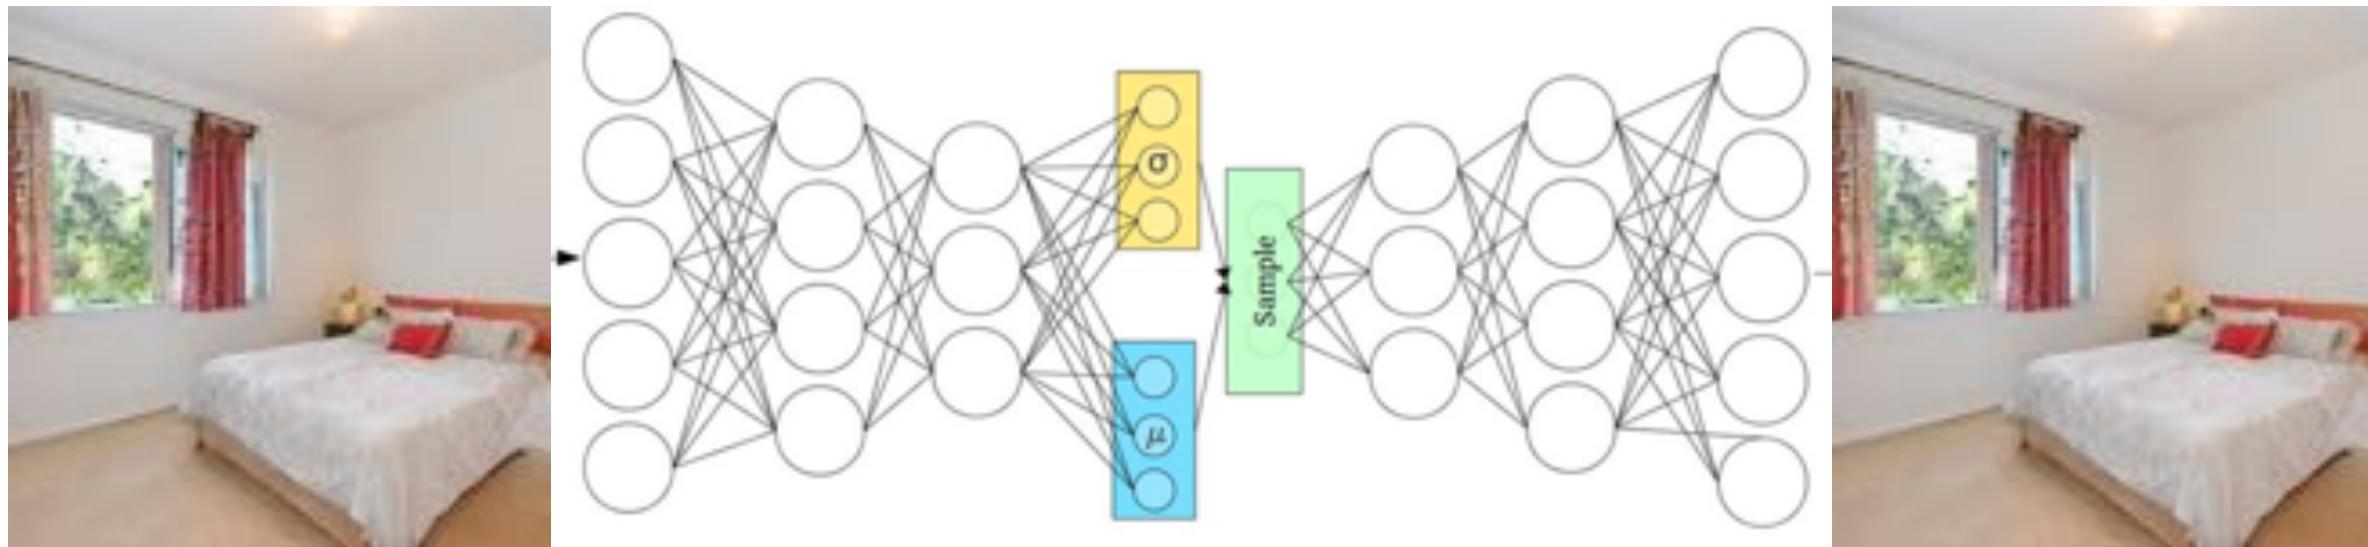


Documents



Variational Autoencoder

- Problem with the vanilla autoencoder is that we cannot sample the latent code to generate new image
- Thus variational autoencoder is introduced



<https://github.com/pytorch/examples/blob/master/vae/main.py>

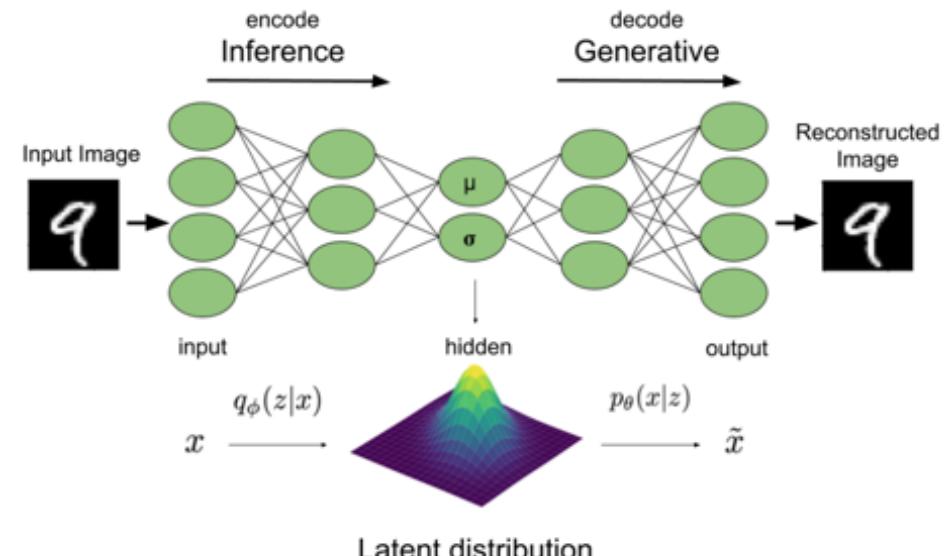
Variational Inference

Evidence Lower Bound (ELBO) holds for any q

$$\begin{aligned}\log p(\mathbf{x}; \theta) &\geq \sum_{\mathbf{z}} q(\mathbf{z}) \log \left(\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right) \\&= \sum_{\mathbf{z}} q(\mathbf{z}) \log p_{\theta}(\mathbf{x}, \mathbf{z}) - \underbrace{\sum_{\mathbf{z}} q(\mathbf{z}) \log q(\mathbf{z})}_{\text{Entropy } H(q) \text{ of } q} \\&= \sum_{\mathbf{z}} q(\mathbf{z}) \log p_{\theta}(\mathbf{x}, \mathbf{z}) + H(q) \\&= \mathcal{L}(\mathbf{x}; \theta, \phi) + D_{KL}(q(\mathbf{z}; \phi) \| p(\mathbf{z}|\mathbf{x}; \theta))\end{aligned}$$

The better $q(\mathbf{z}; \phi)$ can approximate the posterior $p(\mathbf{z}|\mathbf{x}; \theta)$, the smaller $D_{KL}(q(\mathbf{z}; \phi) \| p(\mathbf{z}|\mathbf{x}; \theta))$ we can achieve, the closer ELBO will be to $\log p(\mathbf{x}; \theta)$. Next: jointly optimize over θ and ϕ to maximize the ELBO over a dataset

Variational Inference



$$\begin{aligned}\mathcal{L}(\mathbf{x}; \theta, \phi) &= E_{q_\phi(\mathbf{z}|\mathbf{x})}[\log p(\mathbf{z}, \mathbf{x}; \theta) - \log q_\phi(\mathbf{z}|\mathbf{x})] \\ &= E_{q_\phi(\mathbf{z}|\mathbf{x})}[\log p(\mathbf{z}, \mathbf{x}; \theta) - \log p(\mathbf{z}) + \log p(\mathbf{z}) - \log q_\phi(\mathbf{z}|\mathbf{x})] \\ &= E_{q_\phi(\mathbf{z}|\mathbf{x})}[\log p(\mathbf{x}|\mathbf{z}; \theta)] - D_{KL}(q_\phi(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))\end{aligned}$$

- ① Take a data point \mathbf{x}^i
- ② Map it to $\hat{\mathbf{z}}$ by sampling from $q_\phi(\mathbf{z}|\mathbf{x}^i)$ (*encoder*)
- ③ Reconstruct $\hat{\mathbf{x}}$ by sampling from $p(\mathbf{x}|\hat{\mathbf{z}}; \theta)$ (*decoder*)

What does the training objective $\mathcal{L}(\mathbf{x}; \theta, \phi)$ do?

- First term encourages $\hat{\mathbf{x}} \approx \mathbf{x}^i$ (\mathbf{x}^i likely under $p(\mathbf{x}|\hat{\mathbf{z}}; \theta)$)
- Second term encourages $\hat{\mathbf{z}}$ to be likely under the prior $p(\mathbf{z})$

https://deepgenerativemodels.github.io/assets/slides/cs236_lecture5.pdf

https://deepgenerativemodels.github.io/assets/slides/cs236_lecture6.pdf

Variational Autoencoder

Synthesis quality looked good, but not good for recent two-year standard



VQ-VAE and VQ-VAE-2

- Vector Quantized Variational AutoEncoder (VQ-VAE)
- Discrete latent representation + PixelCNN for handling posterior collapse

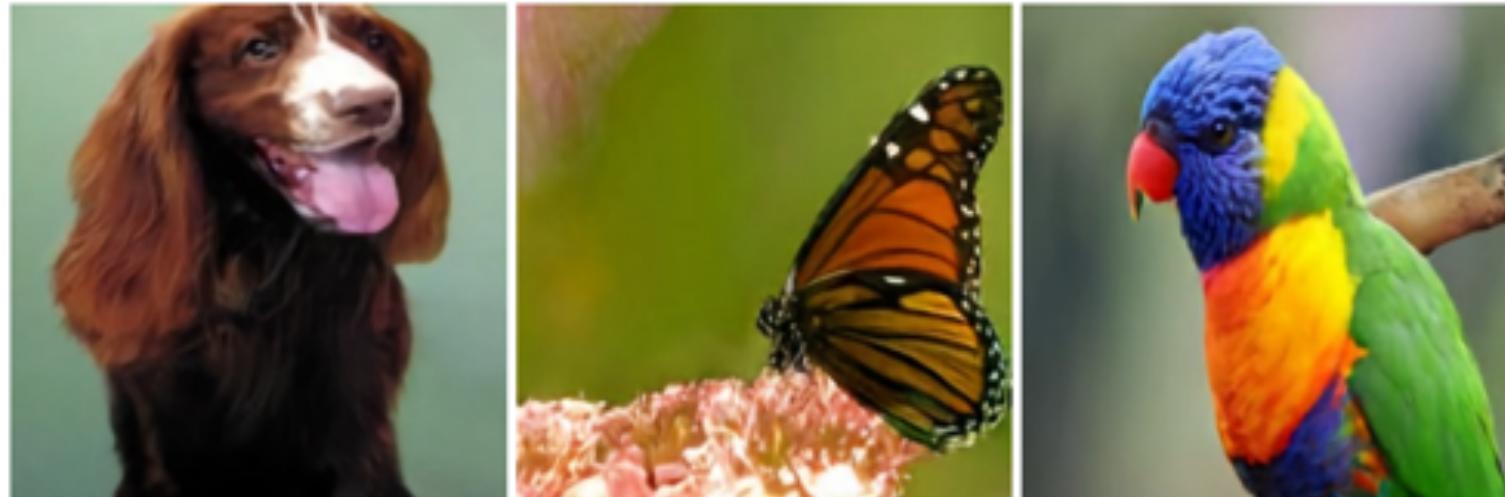
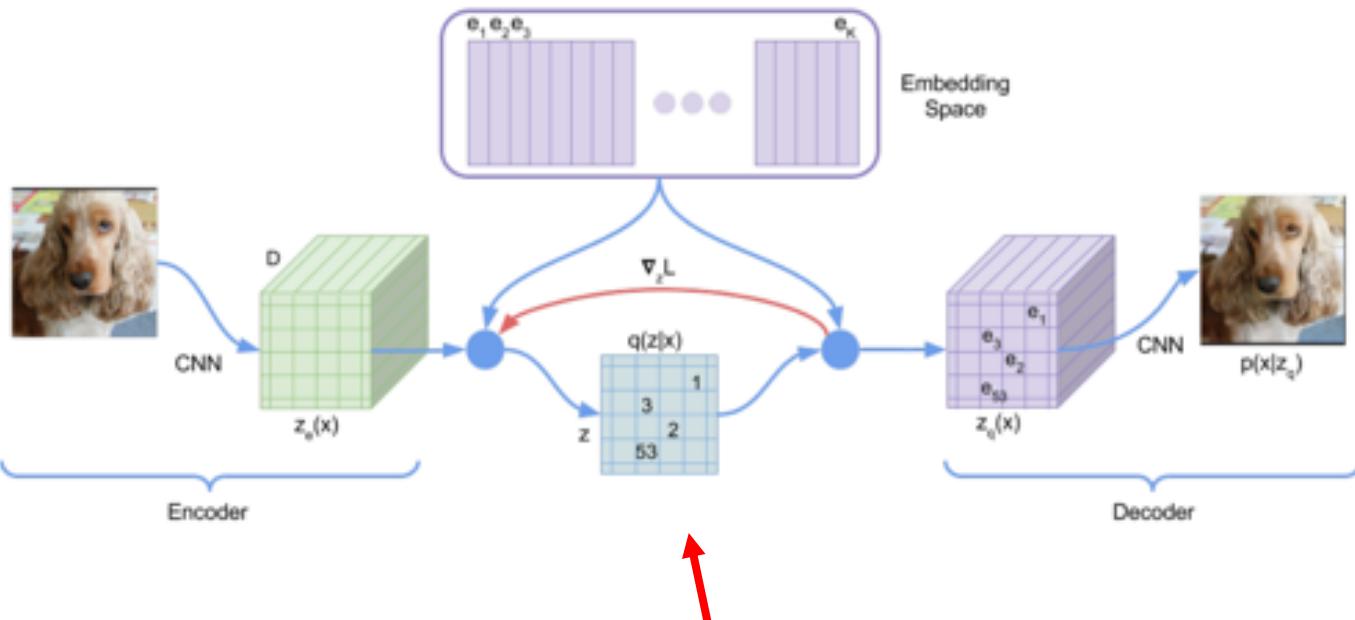


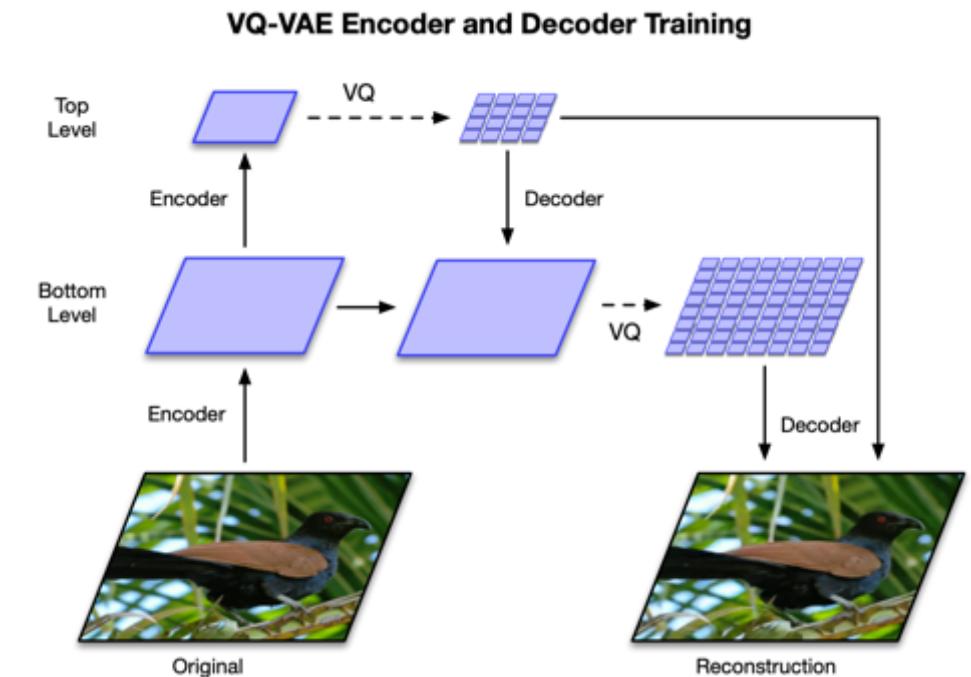
Figure 1: Class-conditional 256x256 image samples from a two-level model trained on ImageNet.

VQ-VAE-1



PixeCNN is learned on the discrete latent space

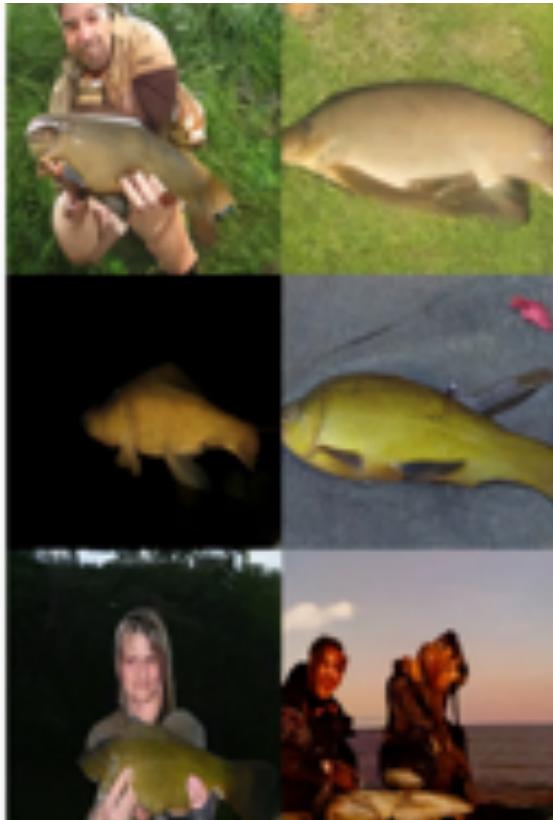
VQ-VAE-2



VQ-VAE and VQ-VAE-2

Very well handle mode-collapse issue in generative models

VQ-VAE on ImageNet's Tinca class



BigGAN

<https://arxiv.org/pdf/1711.00937.pdf>

<https://arxiv.org/pdf/1906.00446.pdf>

Summary of Latent Variable Models

Pros:

- Easy to build flexible models
- Suitable for unsupervised learning

Cons:

- Hard to evaluate likelihood
- Hard to train via maximum-likelihood, as posterior inference $p(z|x)$ is hard which requires variational approximations

Autoregressive models and Variational Autoencoders

- Both model families are based on maximizing likelihoods (or approximation)

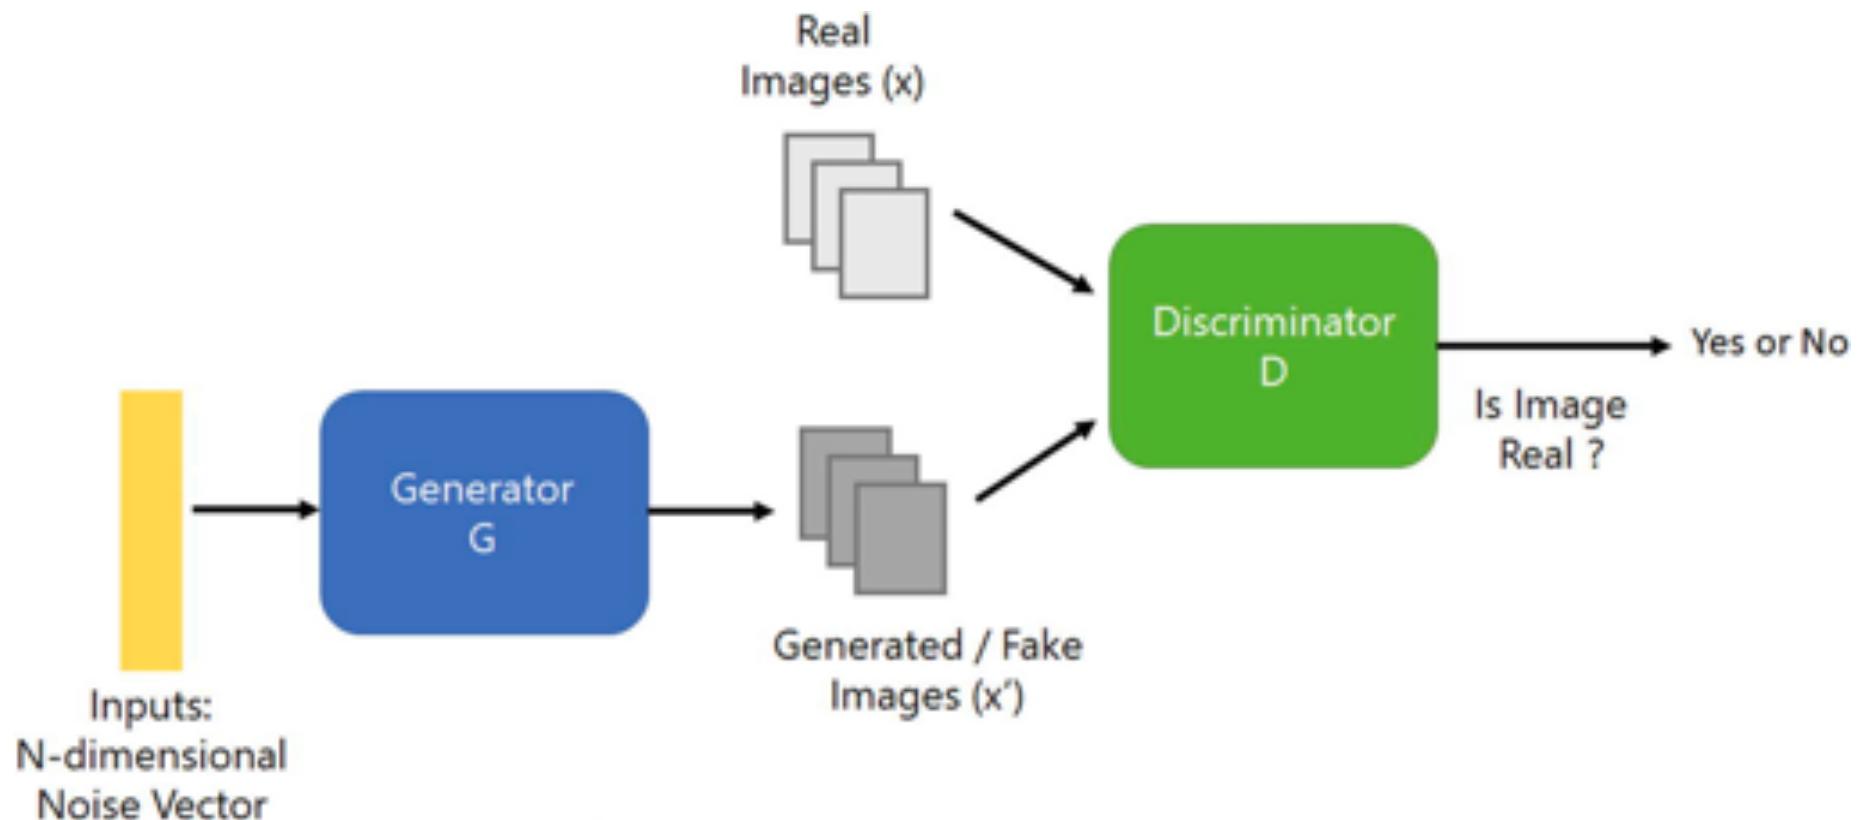
Autoregressive Models: $p_\theta(\mathbf{x}) = \prod_{i=1}^n p_\theta(x_i | \mathbf{x}_{<i})$

Variational Autoencoders: $p_\theta(\mathbf{x}) = \int p_\theta(\mathbf{x}, \mathbf{z}) d\mathbf{z}$

- Is that likelihood a good indicator? Any other likelihood-free generative modeling?
- **Can we throw away KL-divergence and likelihood???**

Generative Adversarial Networks

A two-player minimax game between a generator and a discriminator



Training objective for discriminator:

$$\max_D V(G, D) = E_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + E_{\mathbf{x} \sim p_G} [\log(1 - D(\mathbf{x}))]$$

For a fixed generator G , the discriminator is performing binary classification with the cross entropy objective

- Assign probability 1 to true data points $\mathbf{x} \sim p_{\text{data}}$
- Assign probability 0 to fake samples $\mathbf{x} \sim p_G$

Optimal discriminator

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

Training objective for generator:

$$\min_G V(G, D) = E_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + E_{\mathbf{x} \sim p_G} [\log(1 - D(\mathbf{x}))]$$

For the optimal discriminator $D_G^*(\cdot)$, we have

$$\begin{aligned} & V(G, D_G^*(\mathbf{x})) \\ &= E_{\mathbf{x} \sim p_{\text{data}}} \left[\log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right] + E_{\mathbf{x} \sim p_G} \left[\log \frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right] \\ &= E_{\mathbf{x} \sim p_{\text{data}}} \left[\log \frac{p_{\text{data}}(\mathbf{x})}{\frac{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}{2}} \right] + E_{\mathbf{x} \sim p_G} \left[\log \frac{p_G(\mathbf{x})}{\frac{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}{2}} \right] - \log 4 \\ &= \underbrace{D_{KL} \left[p_{\text{data}}, \frac{p_{\text{data}} + p_G}{2} \right] + D_{KL} \left[p_G, \frac{p_{\text{data}} + p_G}{2} \right]}_{2 \times \text{Jenson-Shannon Divergence (JSD)}} - \log 4 \\ &= 2D_{JSD}[p_{\text{data}}, p_G] - \log 4 \end{aligned}$$

Symmetric KL divergence

Optimal generator
 $p_G = p_{\text{data}}$

Joint loss:

$$\min_{\theta} \max_{\phi} V(G_{\theta}, D_{\phi}) = E_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

GAN training algorithm

Repeat:

- Sample minibatch of m training points $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}$ from \mathcal{D}
- Sample minibatch of m noise vectors $\mathbf{z}^{(1)}, \mathbf{z}^{(2)}, \dots, \mathbf{z}^{(m)}$ from p_z
- Update the generator parameters θ by stochastic gradient **descent**

$$\nabla_{\theta} V(G_{\theta}, D_{\phi}) = \frac{1}{m} \nabla_{\theta} \sum_{i=1}^m \log(1 - D_{\phi}(G_{\theta}(\mathbf{z}^{(i)})))$$

- Update the discriminator parameters ϕ by stochastic gradient **ascent**

$$\nabla_{\phi} V(G_{\theta}, D_{\phi}) = \frac{1}{m} \nabla_{\phi} \sum_{i=1}^m [\log D_{\phi}(\mathbf{x}^{(i)}) + \log(1 - D_{\phi}(G_{\theta}(\mathbf{z}^{(i)})))]$$

GAN Progress for Image Synthesis

2014



GANs [Goodfellow et al.]

2015



DCGAN [Radford et al.]

2017



PG-GAN [Karras et al.]

2018



BigGAN [Brock et al.]

...

GAN Progress for Image Synthesis

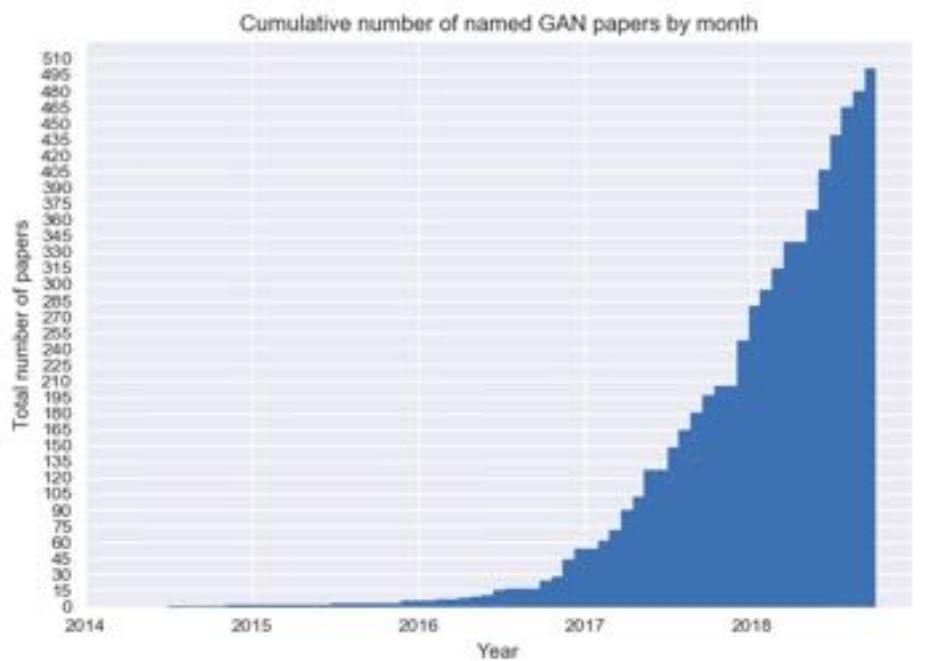


PG-GAN

StyleGAN

BigGAN

GANs explode!



<https://github.com/hindupuravinash/the-gan-zoo>

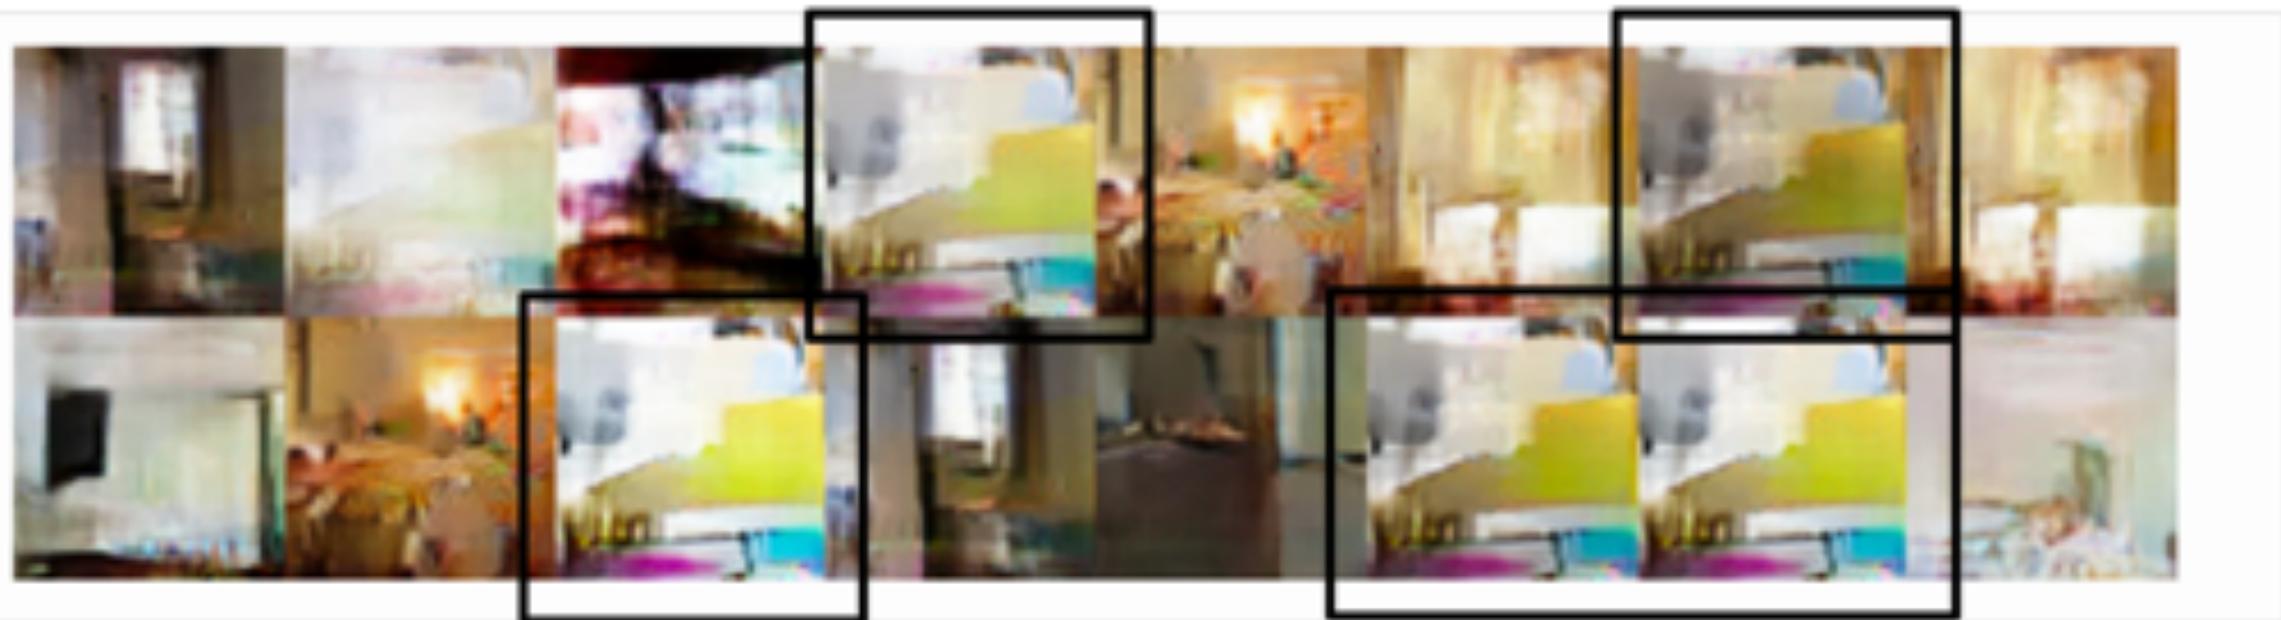
- 3D-ED-GAN
- ARIGAN - [ARIG](#)
- CatGAN - [CatG](#)
- EBGAN - [En](#)
- GANosaic - [I](#)
- 3D-GAN - [L](#)
- ArtGAN - [ArtGA](#)
- CausalGAN - [C](#)
- ecGAN - [eC](#)
- GANVO - [GA](#)
- 3D-IWGAN
- ASDL-GAN - [At](#)
- CC-GAN - [Sen](#)
- ED//GAN - [E](#)
- Networks
- 3D-PhysNet
- ATA-GAN - [Atte](#)
- cd-GAN - [Con](#)
- Editable GAI
- GAP - [Conte](#)
- 3D-RecGAN
- Attention-GAN
- CDcGAN - [Sir](#)
- EGAN - [Enh](#)
- GAP - [Gener](#)
- ABC-GAN -
- AttGAN - [Arbitr](#)
- CE-GAN - [Dee](#)
- EL-GAN - [El](#)
- GATS - [Sam](#)
- ABC-GAN -
- AttnGAN - [Attn](#)
- CFG-GAN - [Co](#)
- ELEGANT -
- AC-GAN - [C](#)
- AVID - [AVID: Ad](#)
- CGAN - [Conditi](#)
- EnergyWGA
- acGAN - [Fac](#)
- B-DCGAN - [B-D](#)
- CGAN - [Contr](#)
- ESRGAN - [E](#)
- ACGAN - [Cc](#)
- b-GAN - [Gener](#)
- Chekhov GAN
- ExGAN - [Eye](#)
- acGAN - [On](#)
- BAGAN - [BAGA](#)
- ciGAN - [Conditi](#)
- ExposureGA
- ACTuAL - [AC](#)
- Bayesian GAN -
- CinCGAN - [Un](#)
- ExprGAN - [E](#)
- AdaGAN - [A](#)
- Bayesian GAN -
- CipherGAN - [L](#)
- f-CLSWGAN
- Adaptive GA
- BCGAN - [Bayes](#)
- ClusterGAN - [C](#)
- GLCA-GAN
- AdvEntuRe -
- BCGAN - [Bidire](#)
- CM-GAN - [CM](#)
- GM-GAN - [C](#)
- AdvGAN - [G](#)
- BEAM - [Boltzma](#)
- CoAtt-GAN - [A](#)
- Fairness GAI
- AE-GAN - [A](#)
- BEGAN - [BEGA](#)
- CoGAN - [Coup](#)
- FakeGAN - [F](#)
- AE-OT - [Lat](#)
- BEGAN-CS - [Es](#)
- ComboGAN - [C](#)
- FBGAN - [Fei](#)
- AEGAN - [Le](#)
- Bellman GAN - [I](#)
- ConceptGAN -
- FGNet - [GO](#)
- AF-DCGAN
System
- BGAN - [Binary C](#)
- Conditional cyc
- FC-GAN - [Fi](#)
- Bi-GAN - [Auton](#)
- contrast-GAN
- FF-GAN - [Tc](#)
- AffGAN - [Ar](#)
- Aided Genetic A
- Context-RNN-
- FGGAN - [Ac](#)
- AIM - [Gener](#)
- BicycleGAN - [To](#)
- CorrGAN - [Cor](#)
- GRAN - [Gen](#)
- AL-CGAN -
- BiGAN - [Advers](#)
- Coulomb GAN
- Graphical-G
- ALI - [Advers](#)
- BinGAN - [BinGA](#)
- Cover-GAN - [C](#)
- Fila-GAN - [S](#)
- AlignGAN -
- BourGAN - [Bou](#)
- cowboy - [Defe](#)
- First Order G
- AlphaGAN -
- BranchGAN - [Bi](#)
- CR-GAN - [CR-](#)
- GraspGAN -
- AM-GAN - [A](#)
- BRE - [Improving](#)
- Cramèr GAN -
- HAN - [Chine](#)
- AmbientGAN
- BridgeGAN - [Ge](#)
- Cross-GAN - [C](#)
- HAN - [Bidire](#)
- AMC-GAN -
- BS-GAN - [Boun](#)
- crVAE-GAN - [C](#)
- HiGAN - [Exp](#)
- AnoGAN - [L](#)
- BubGAN - [BubC](#)
- CS-GAN - [Imp](#)
- HP-GAN - [H](#)
- APD - [Adver](#)
- BWGAN - [Banana](#)
- CSG - [Speech](#)
- FusedGAN -
- APE-GAN -
- C-GAN - [Face A](#)
- CT-GAN - [CT-I](#)
- CSG - [Speech](#)
- FusionGAN -
- hredGAN - [N](#)

Challenges for GANs

- Many tricks have been applied to train GANs
- There are still many challenges:
 - Unstable optimization
 - Mode collapse
 - Evaluation: how to evaluate the synthesis result

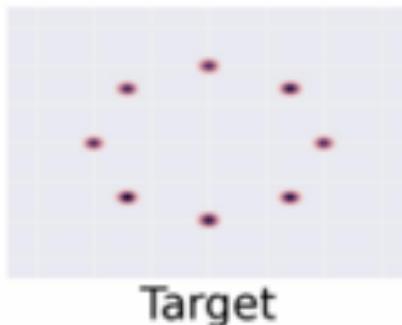
Mode Collapse

GAN often collapses to one or few samples (modes)



Arjovsky et al., 2017

Mode Collapse



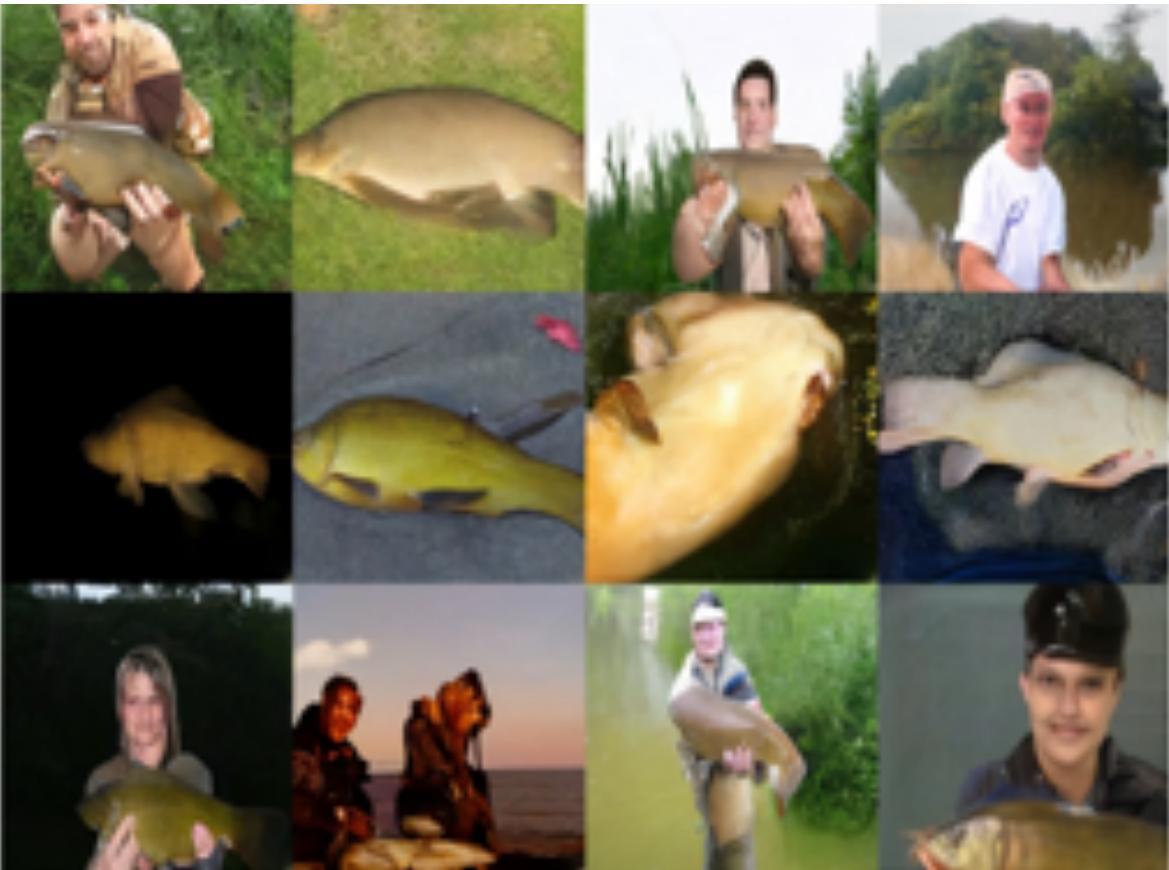
- True distribution is a mixture of Gaussians



Source: Metz et al., 2017

Comparison of VQ-VAE and BigGAN

Likelihood-based



Adversarial learning-based



Comparison of VQ-VAE and BigGAN

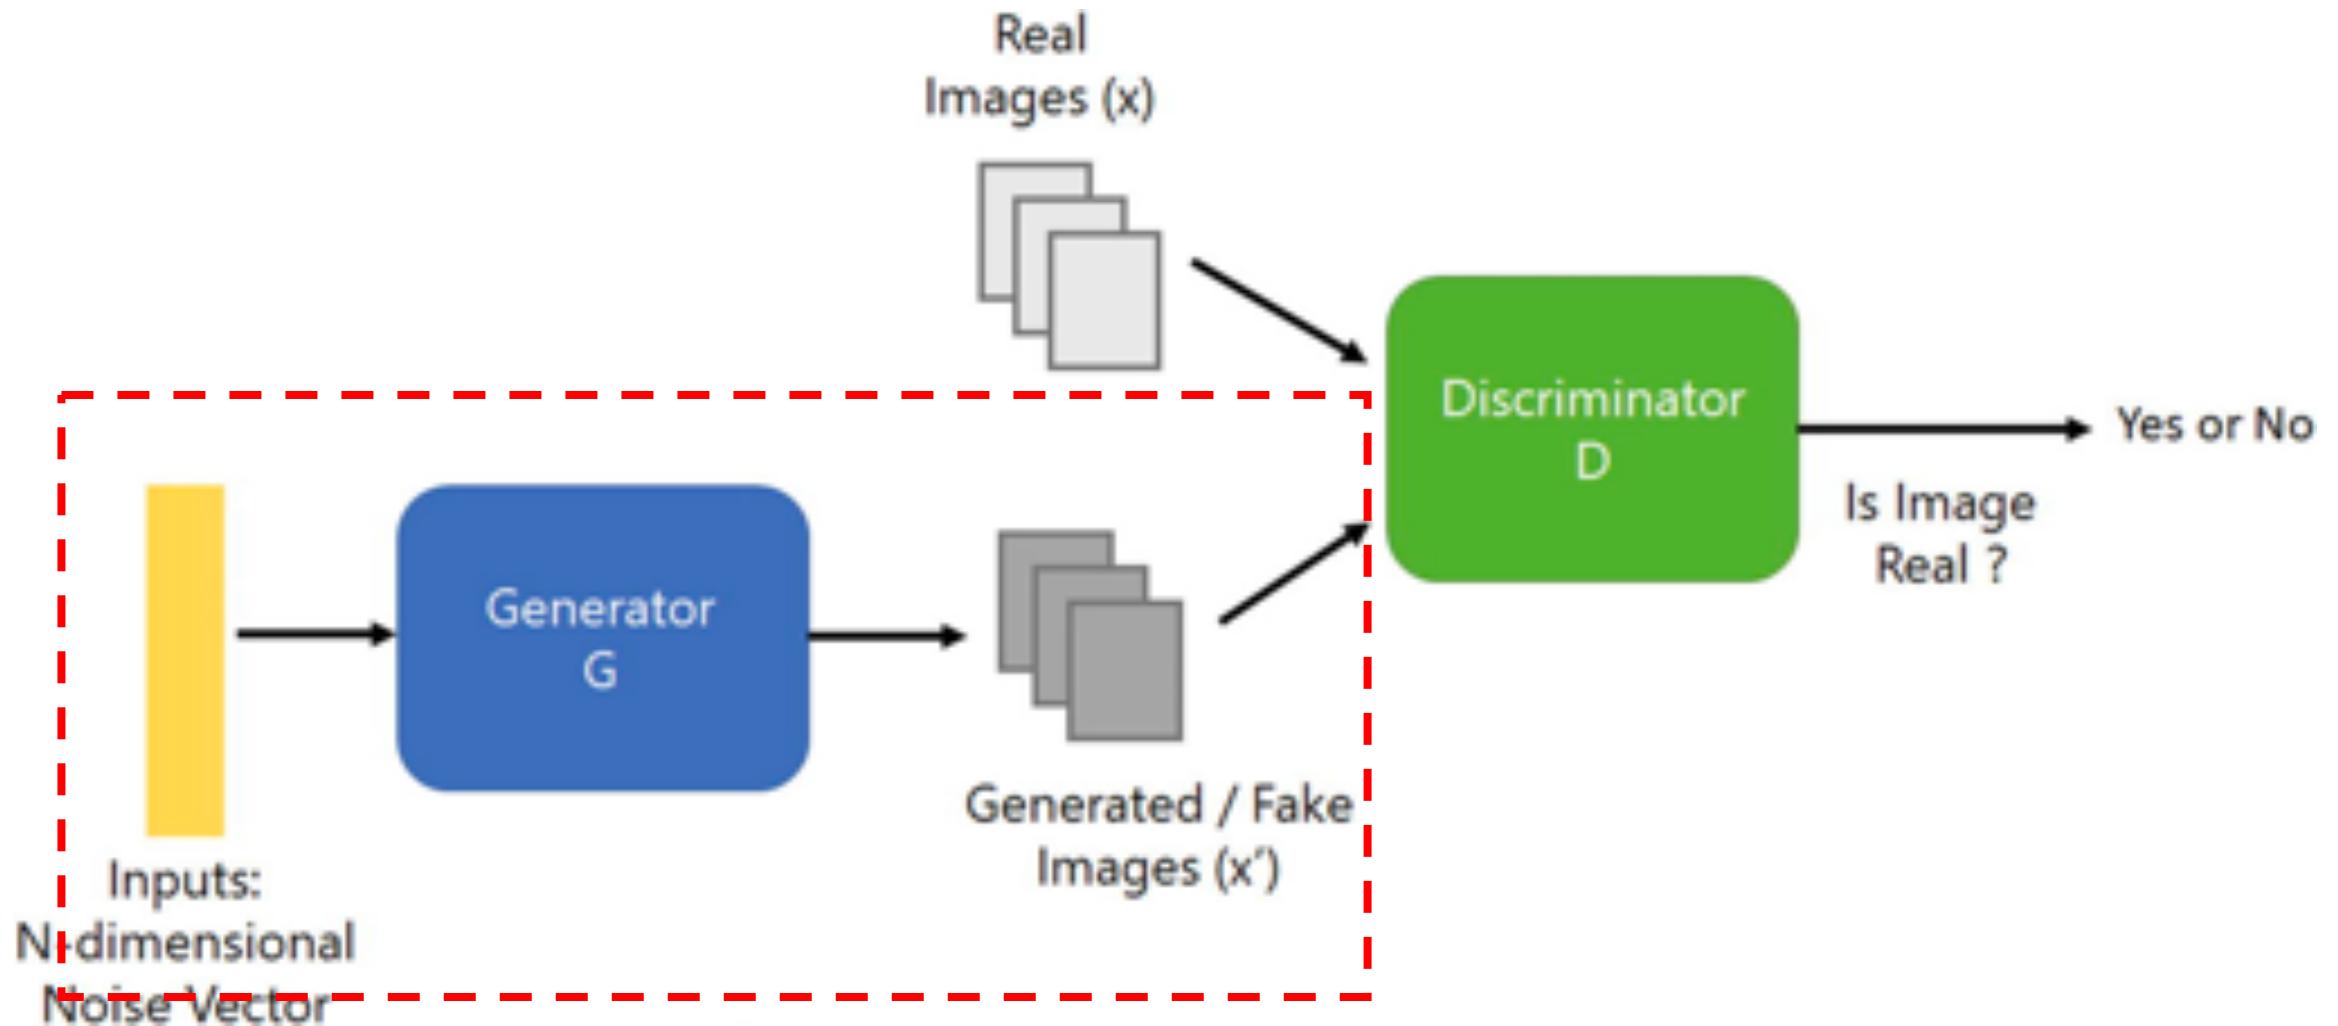
Likelihood-based



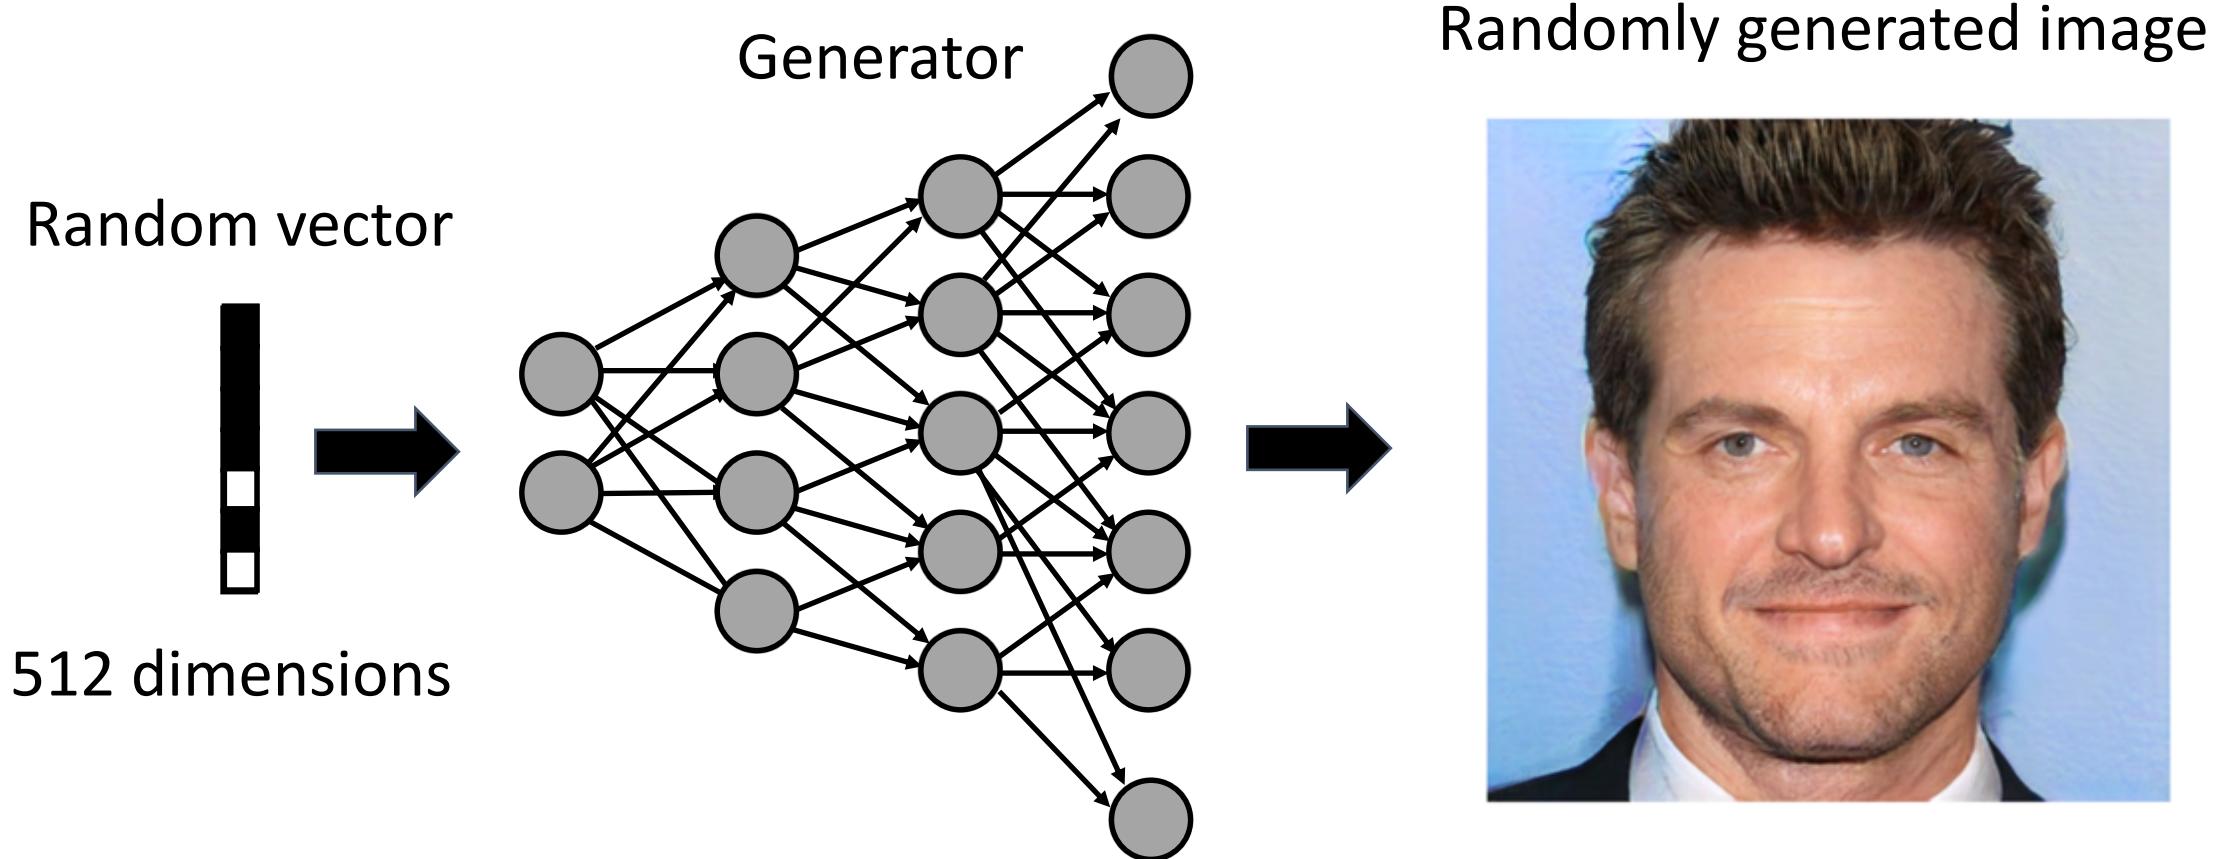
Adversarial learning-based



What have been learned inside the GANs?



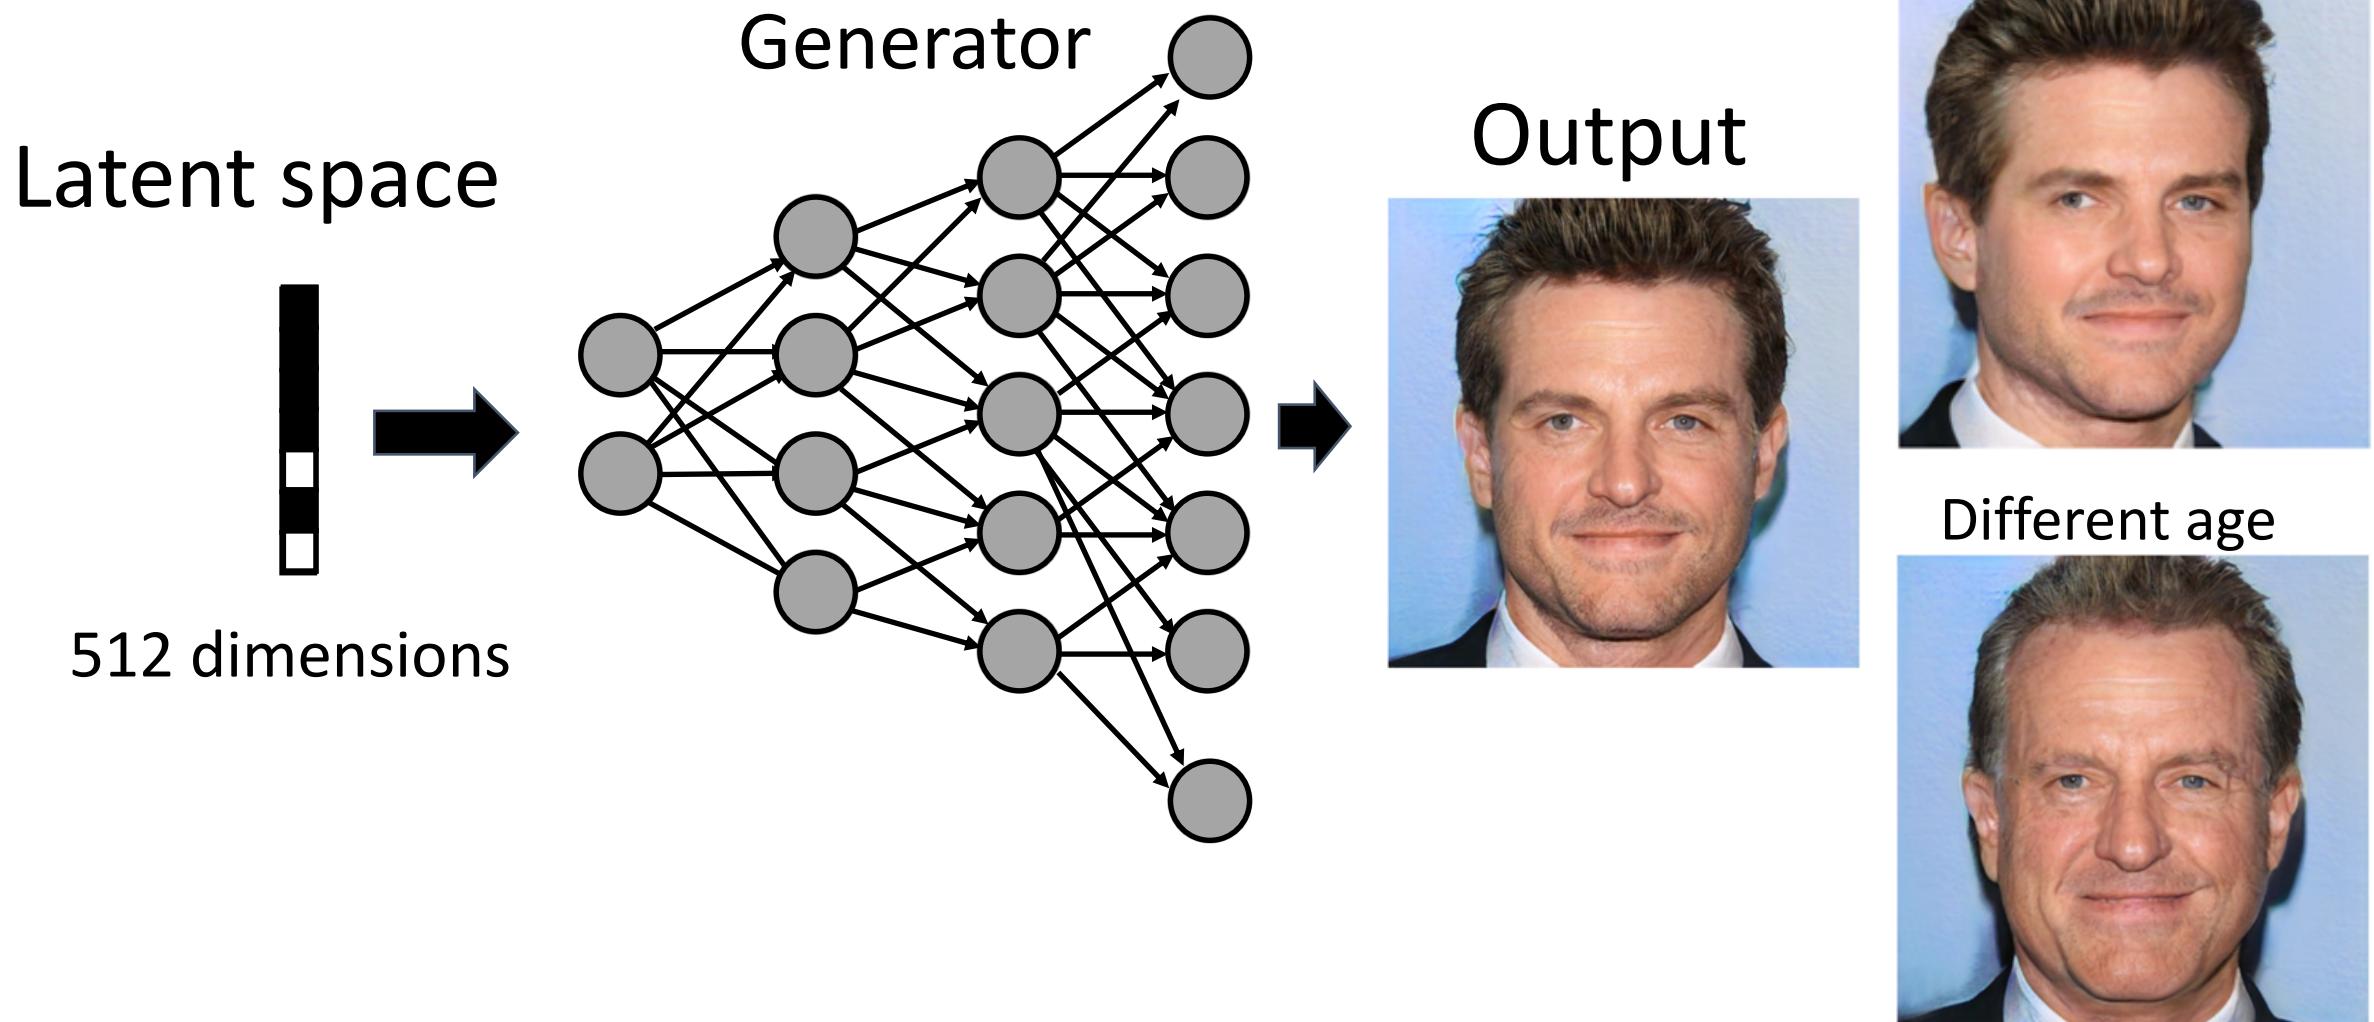
Deep Generative Representation



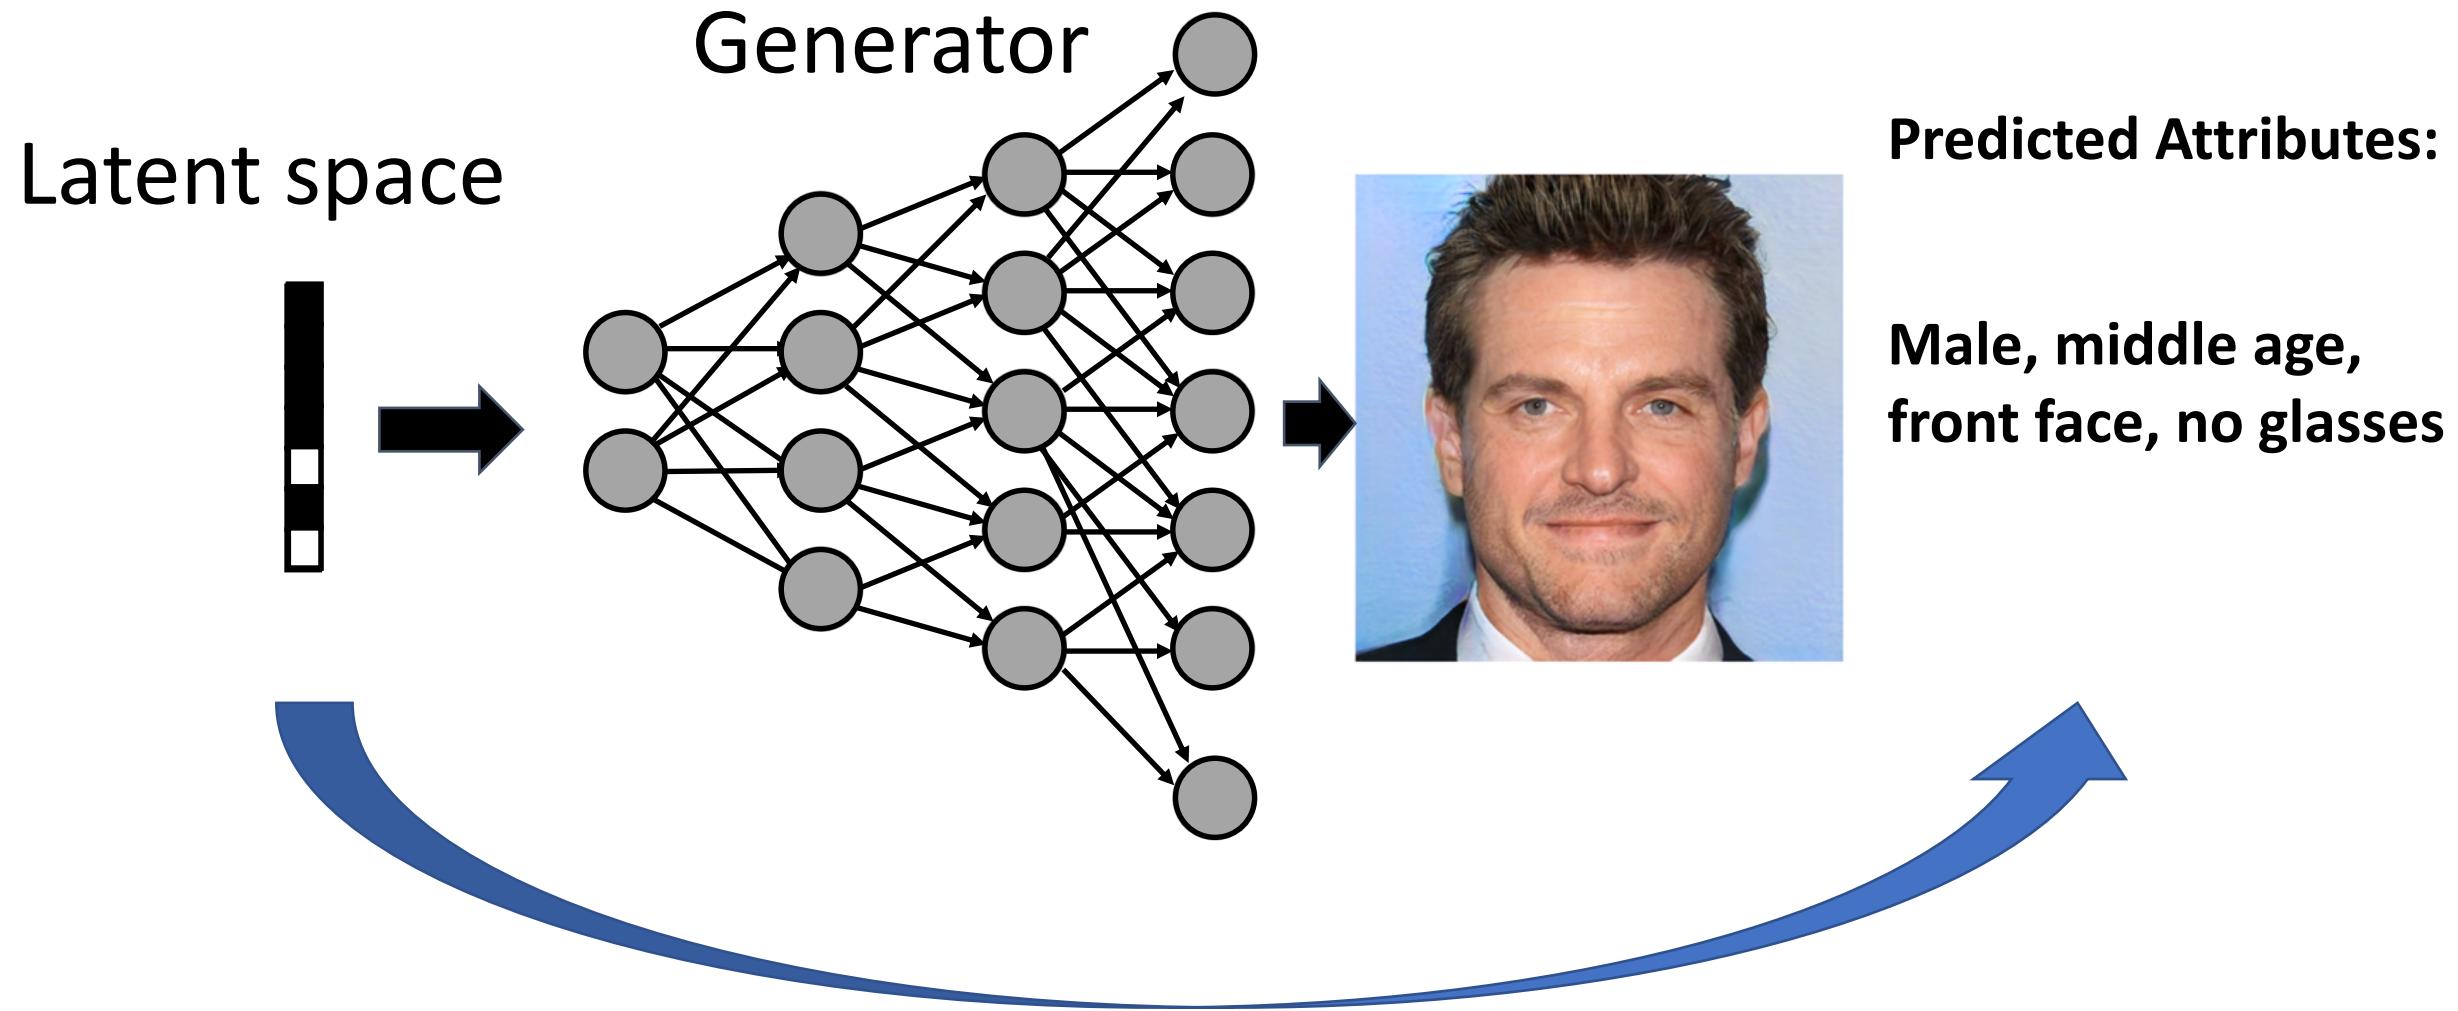
Random Walk in the Latent Space



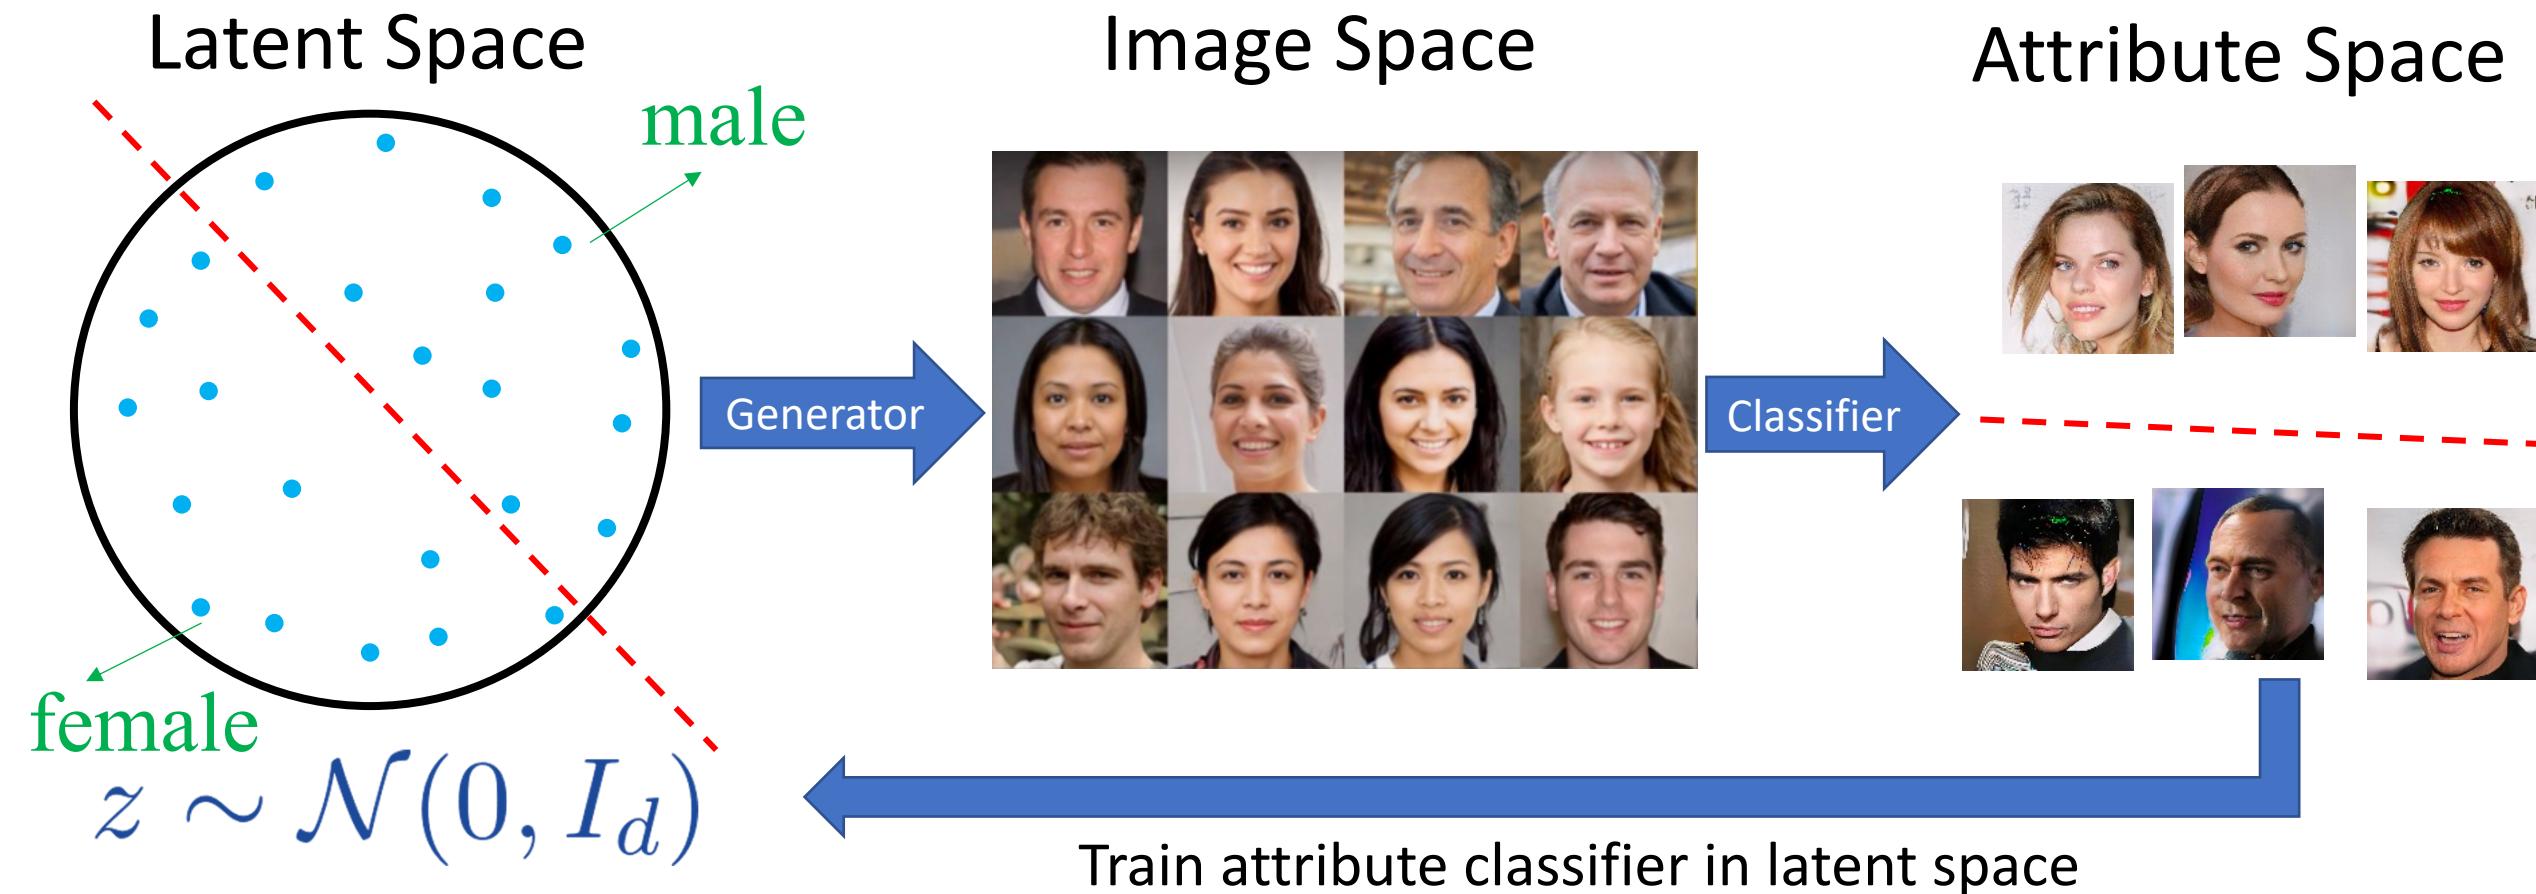
How to Customize the Synthesis Output?



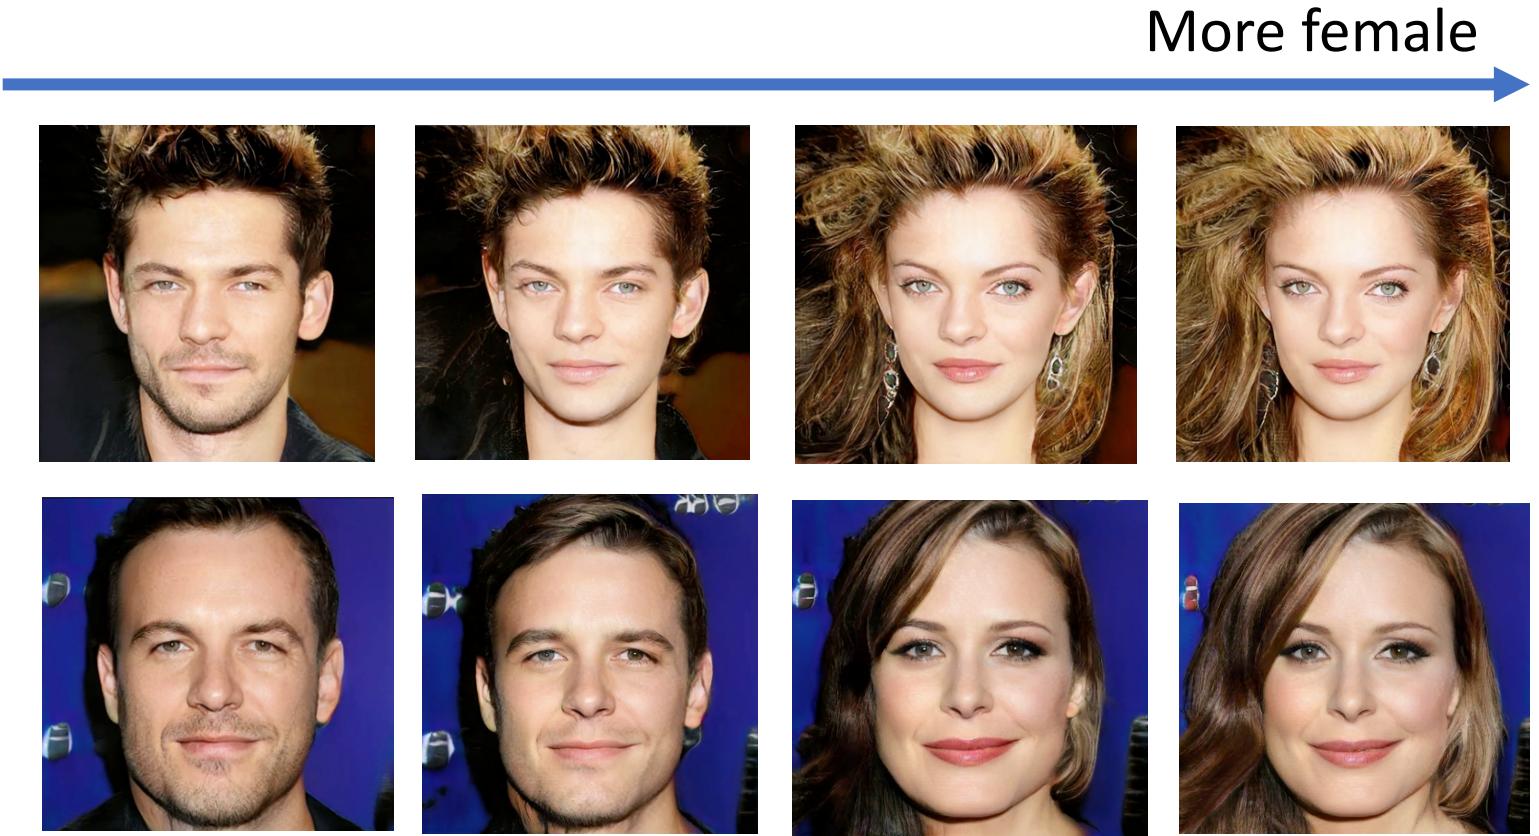
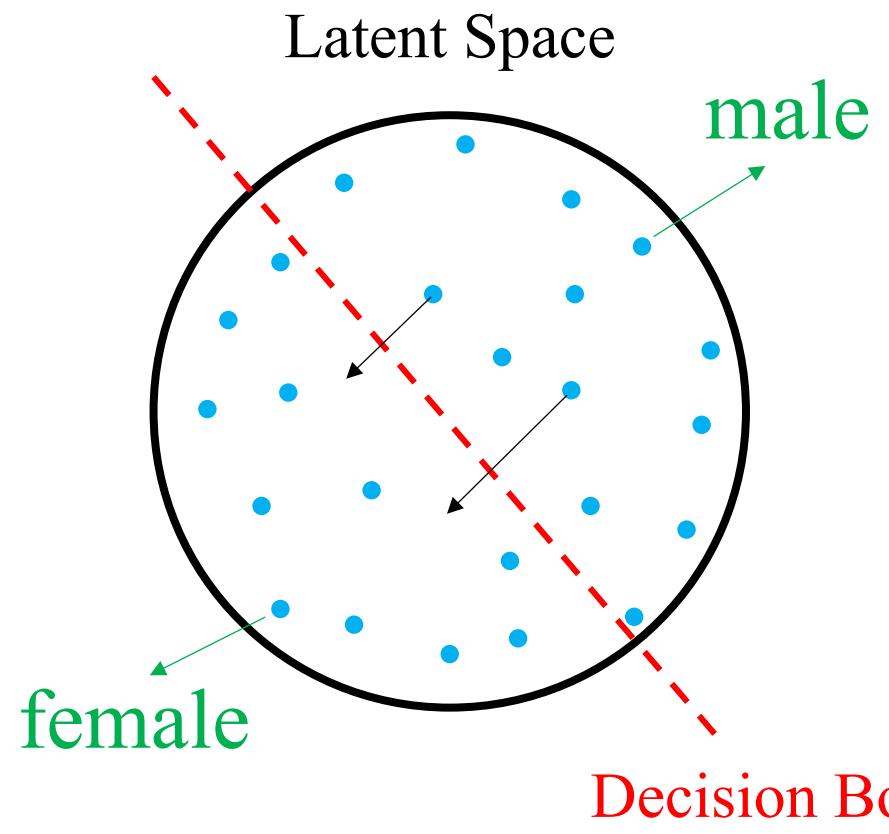
Causal Relations in Latent Space



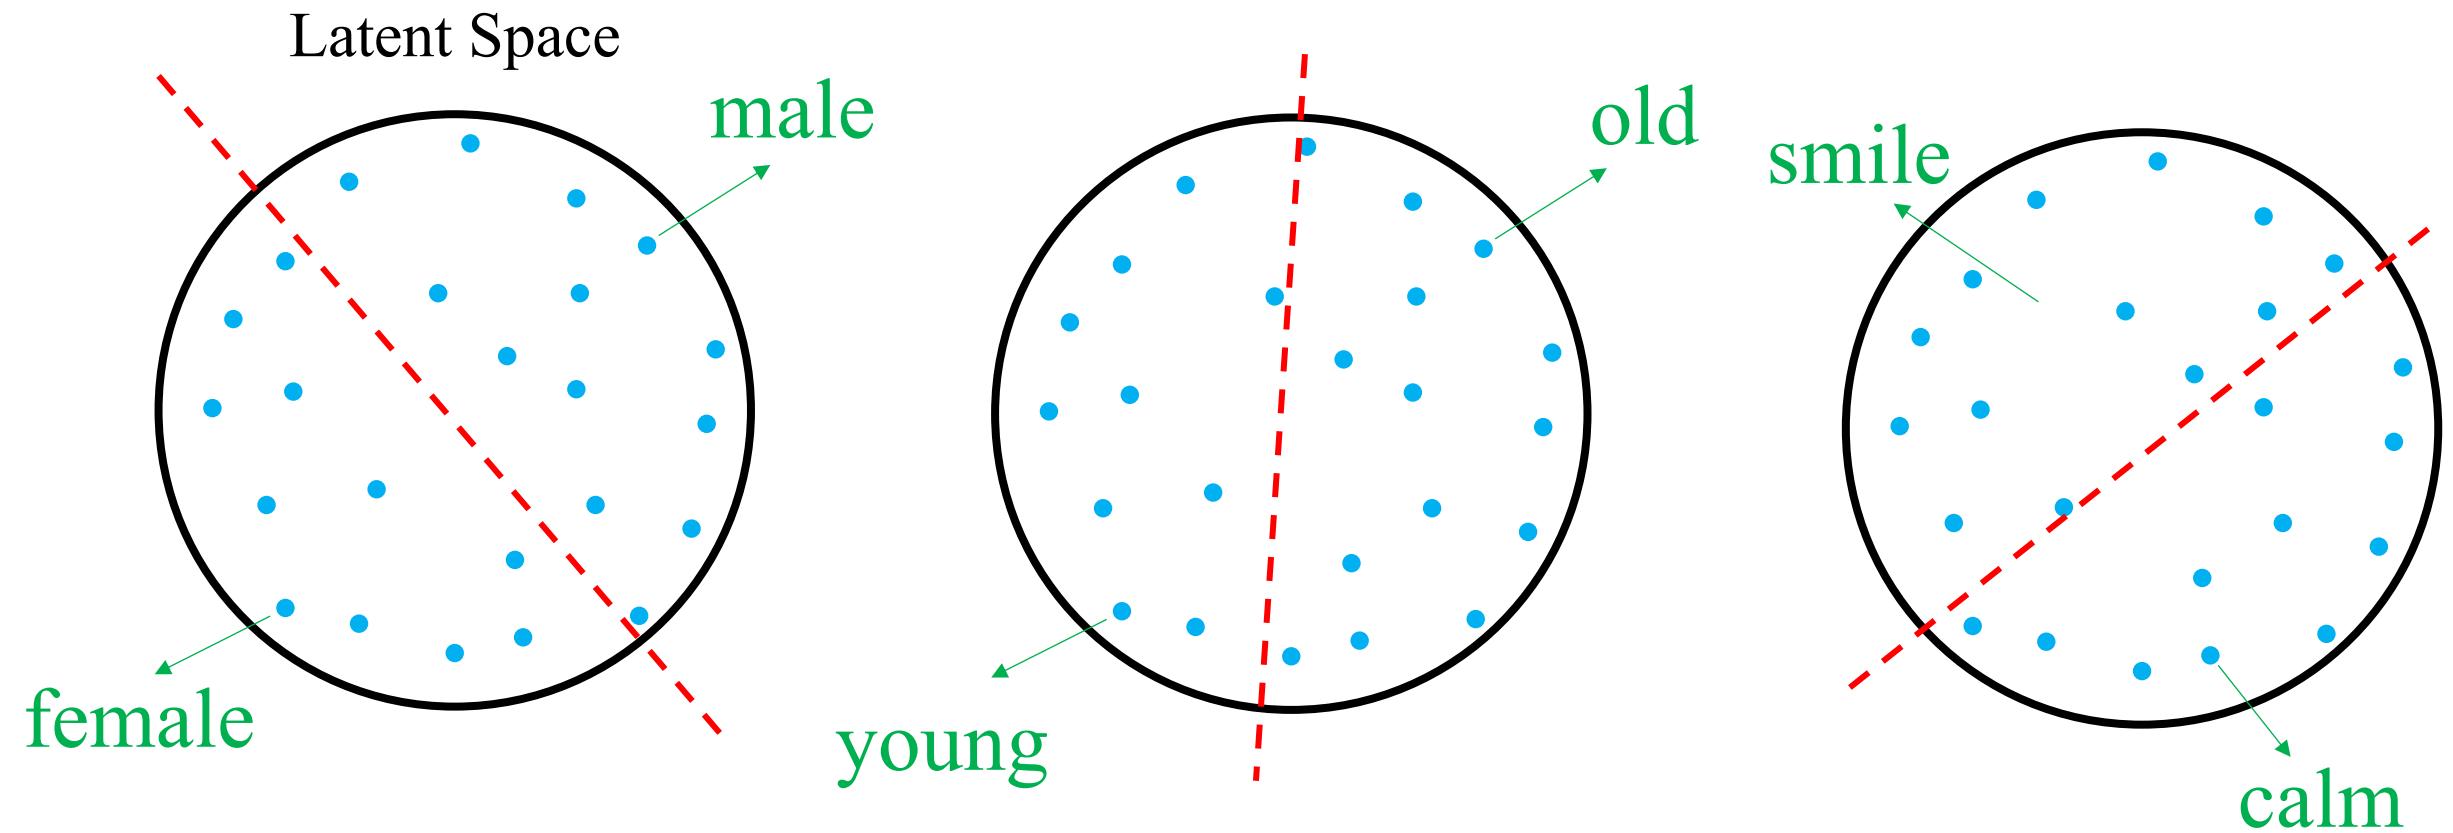
InterFaceGAN: Interpreting Semantics in Face GANs



Varying the Latent Code through Boundary



Various Attribute Boundaries to Divide the Latent Space



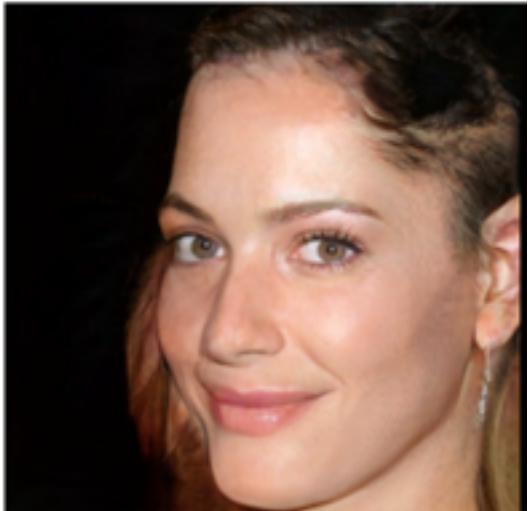
Make me cooler



Make me younger



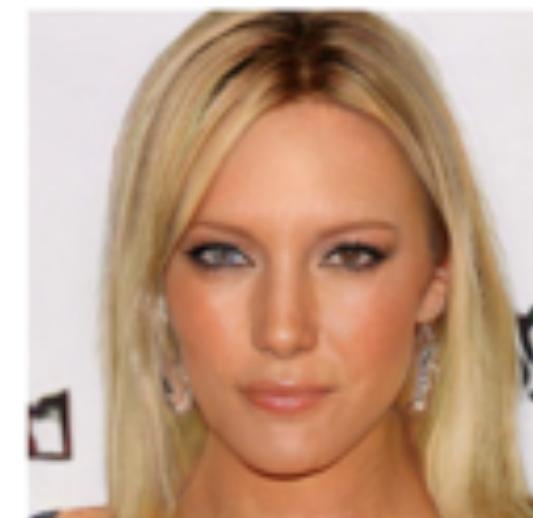
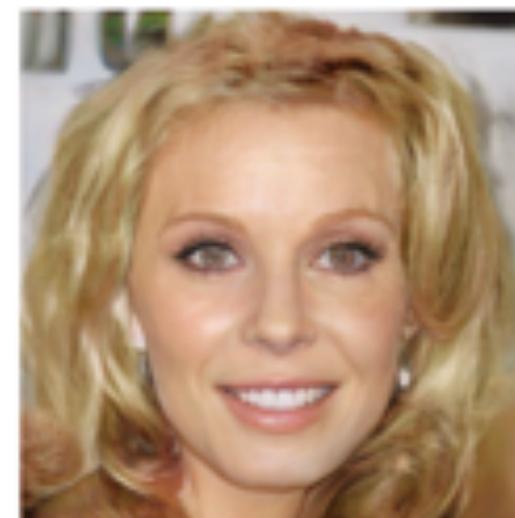
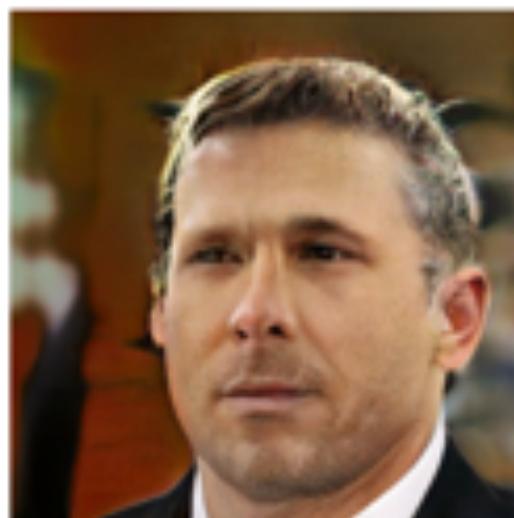
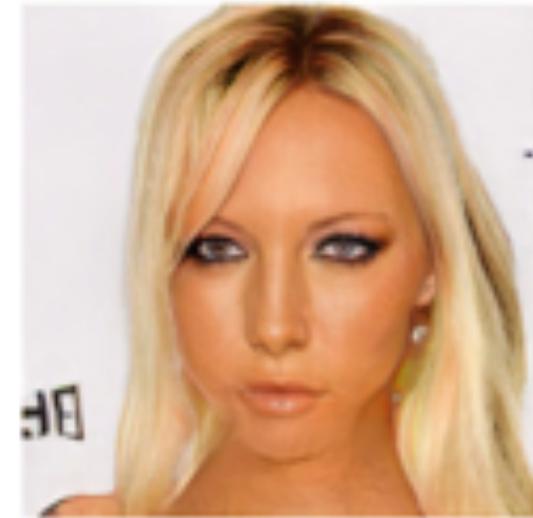
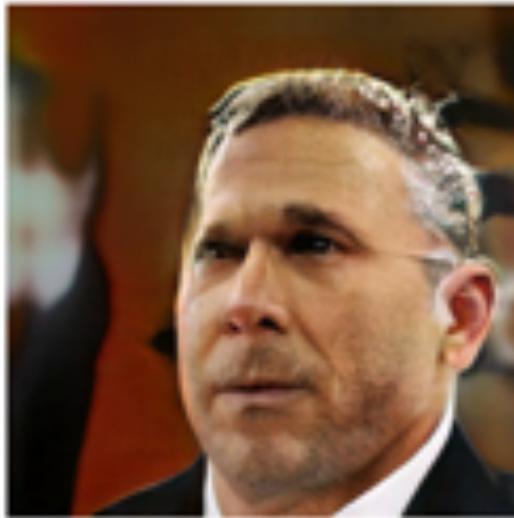
Make me front faced



Make me more man



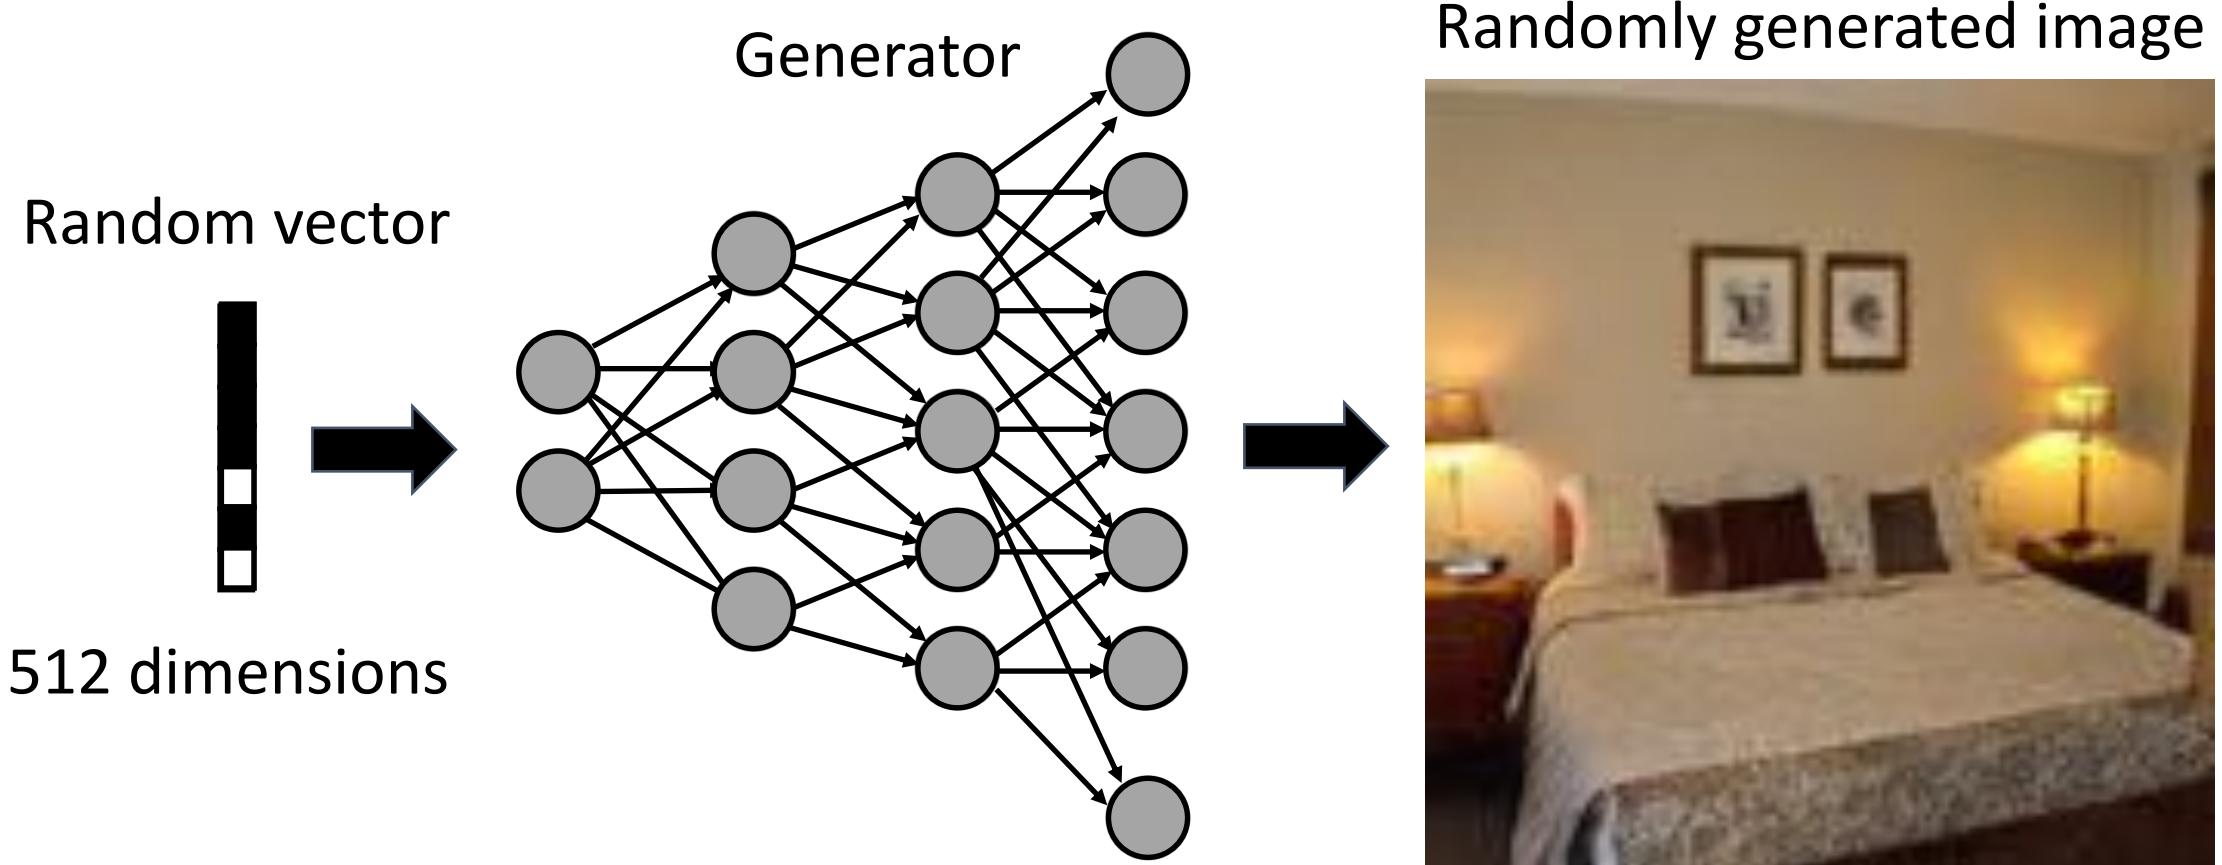
Correcting the Mistakes Made by GAN



Demo Video

Face Manipulation
with InterpretGAN

Latent Semantics in GANs for Scene Synthesis



Random Walk in Latent Space of Bedroom



Multiple Levels of Abstractions for Scenes

Scene category:

bedroom

Scene attributes:

nature lighting

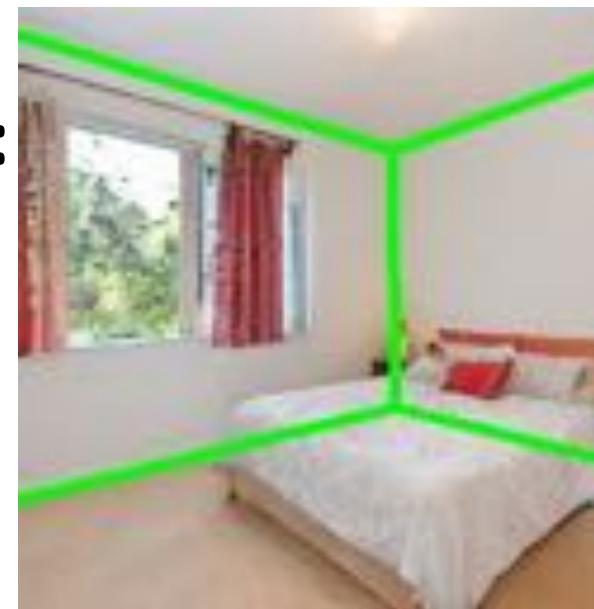
wood

foliage

...



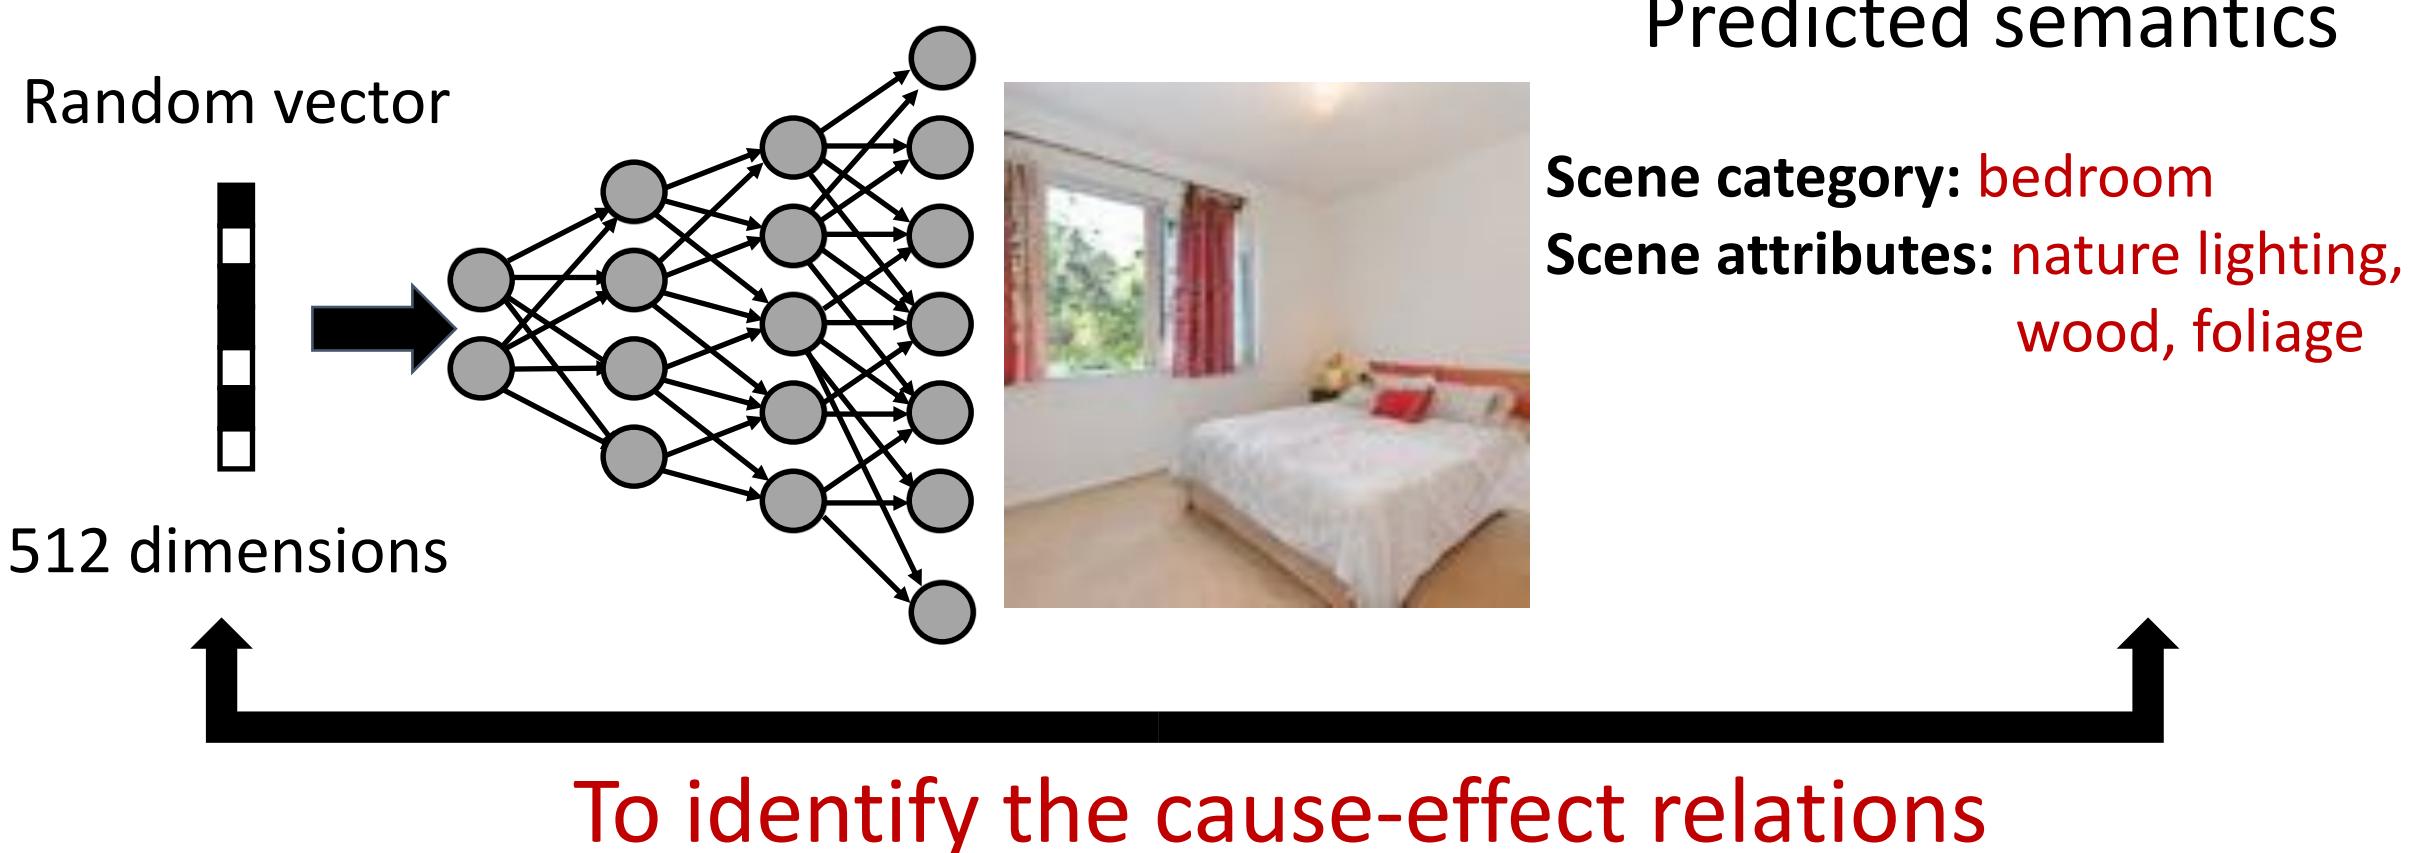
Layout



Segmentation



Identifying the Causality in Latent Space



Result on turning up the lights

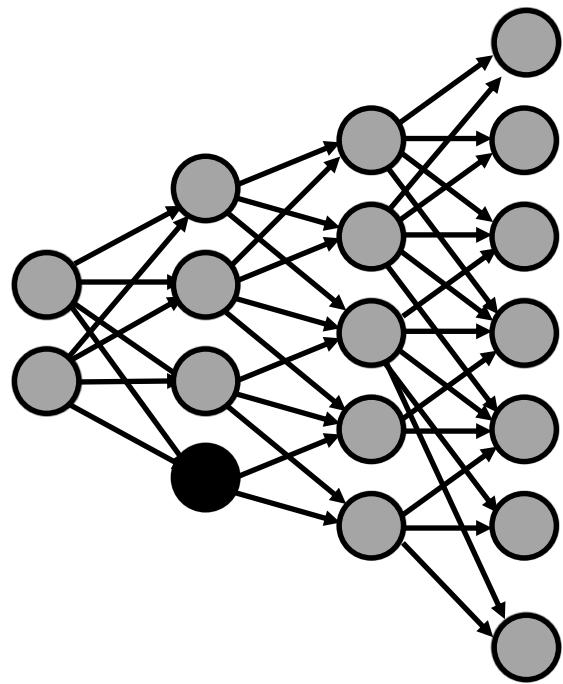
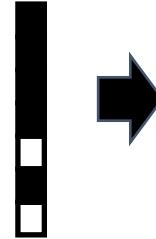


Result on ageing the scenes



Layer-wise Stochasticity

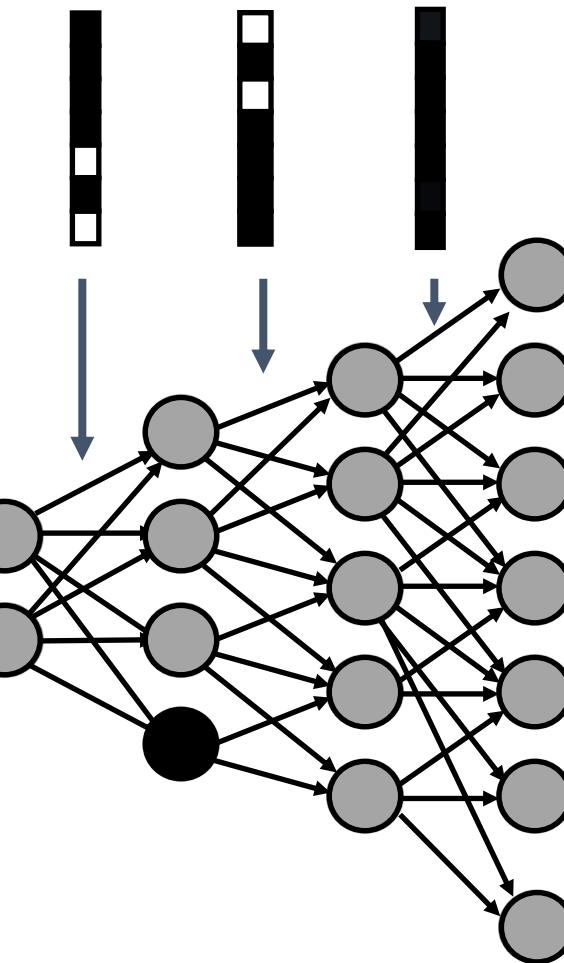
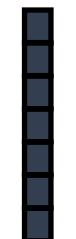
Stochastic
vector



DC-GAN, PG-GAN

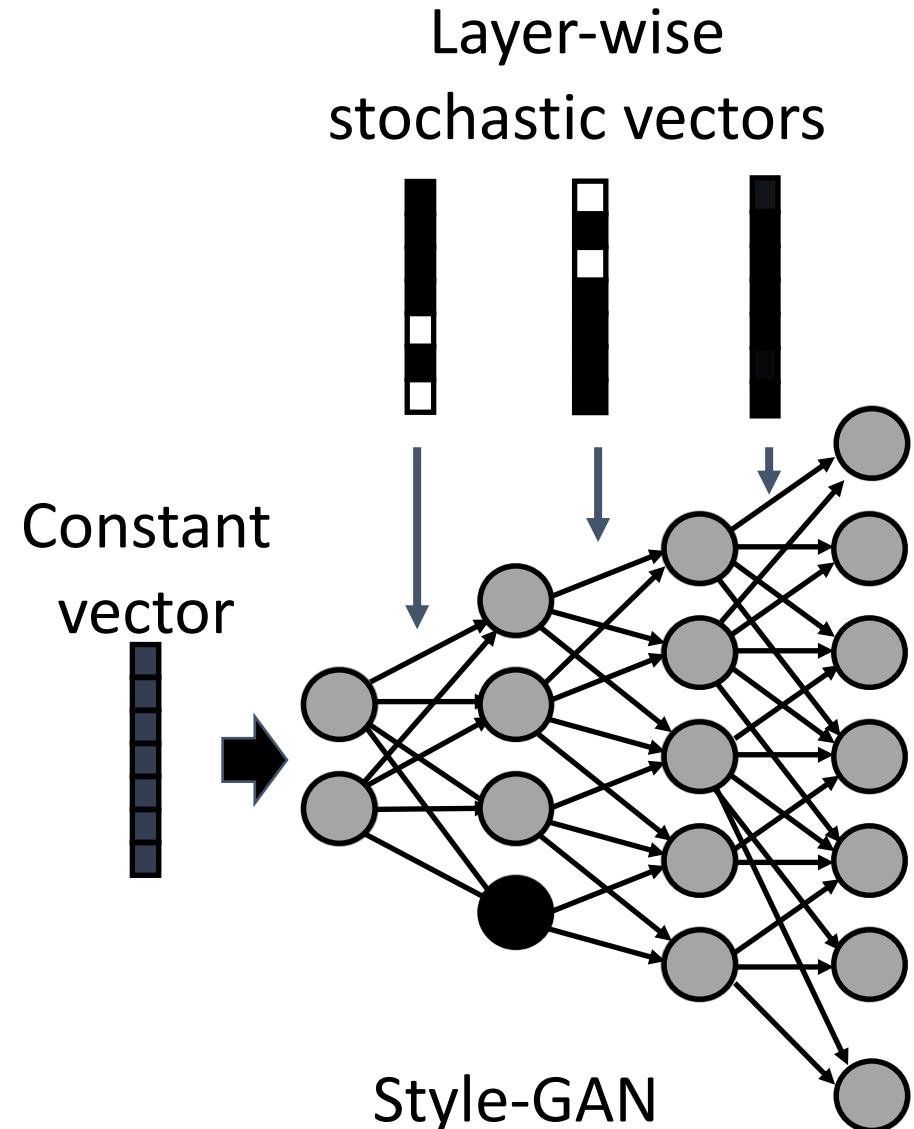
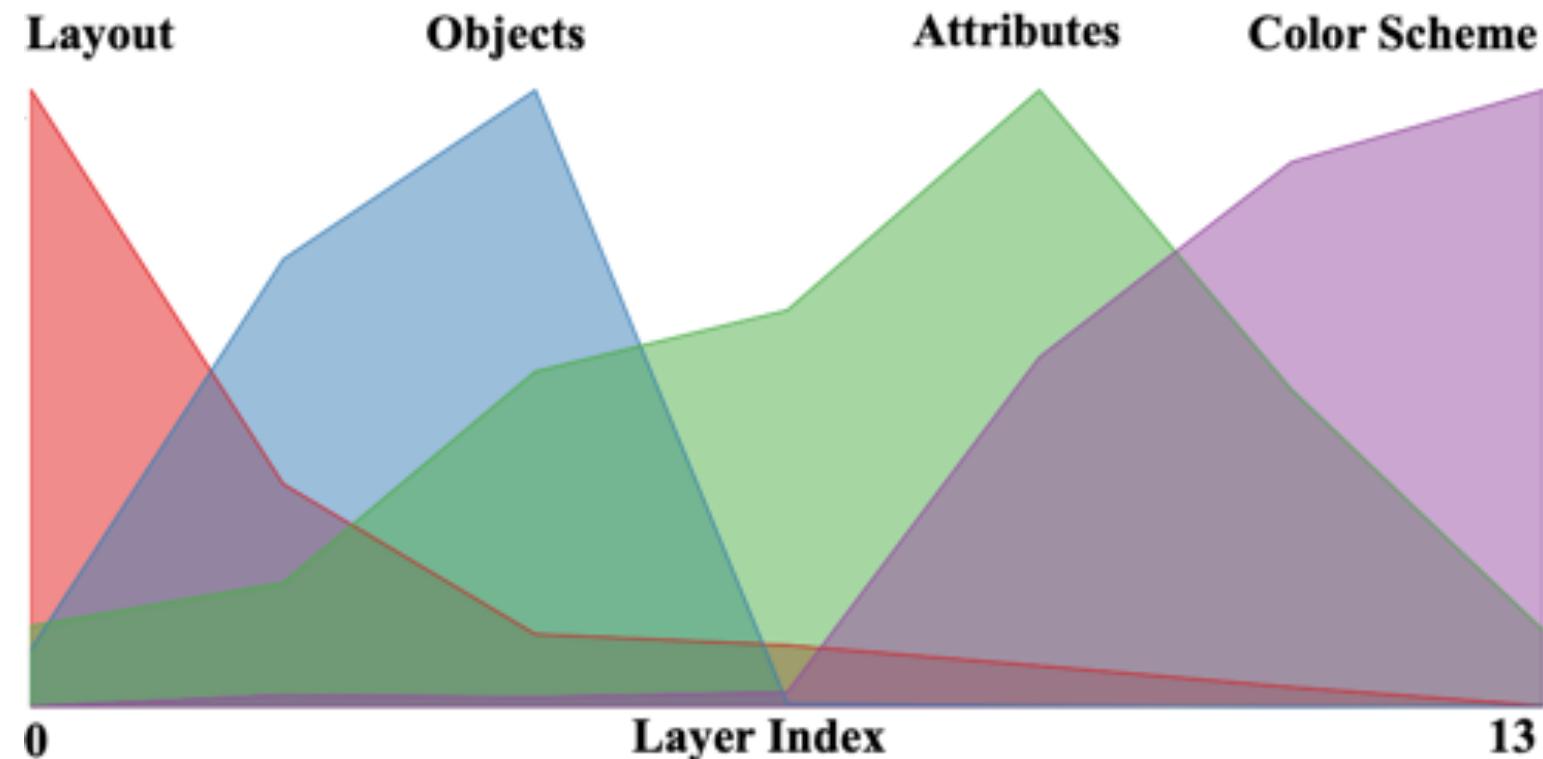
Layer-wise
stochastic vectors

Constant
vector

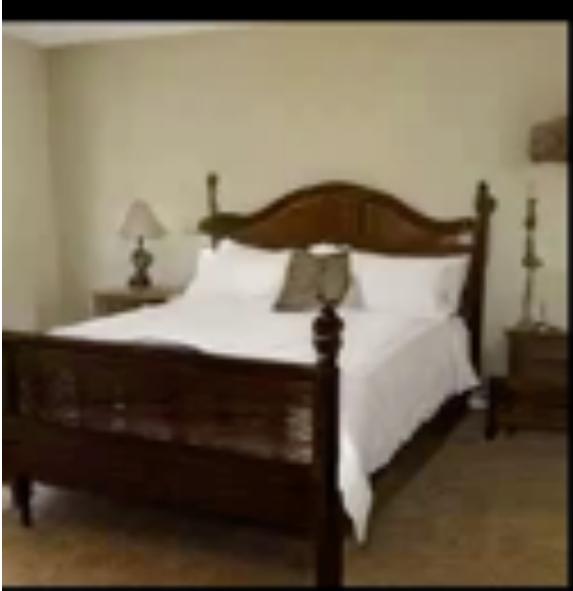


Style-GAN [Karras et al]

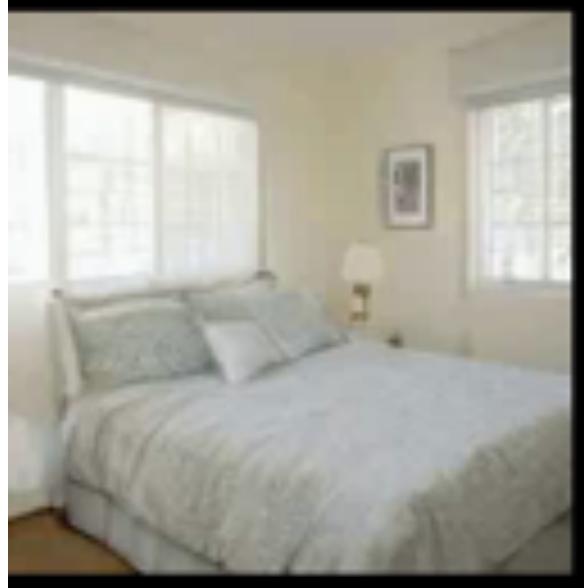
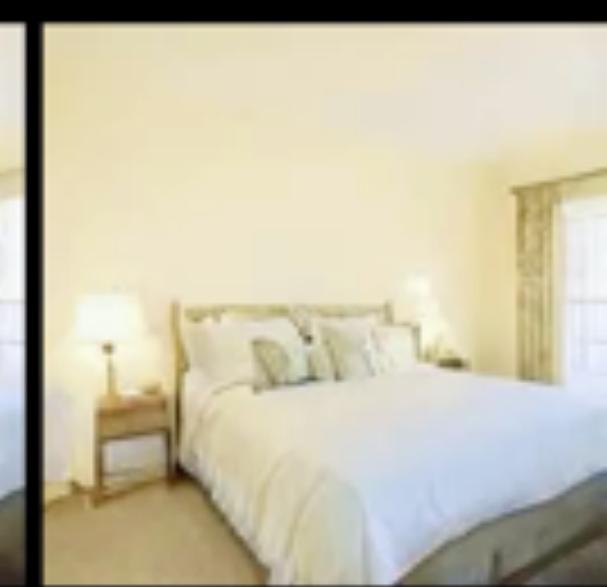
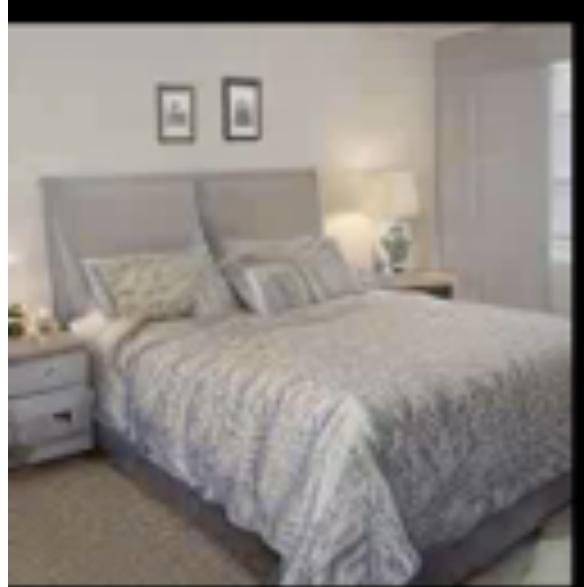
Semantic hierarchy emerges across the layers of generator



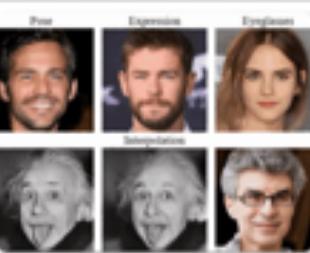
Varying viewpoints



Bedroom to Dining Room



Projects



In-Domain GAN Inversion for Real Image Editing

Jiapeng Zhu*, Yujun Shen*, Deli Zhao, Bolei Zhou
arXiv.2004.00049 preprint

[Paper] [Project Page] [Code] [Demo Video]



Semantic Hierarchy Emerges in Deep Generative Representations for Scene Synthesis

Ceyuan Yang*, Yujun Shen*, Bolei Zhou
arXiv.1911.09267 preprint

[Paper] [Project Page] [Code] [Demo Video]

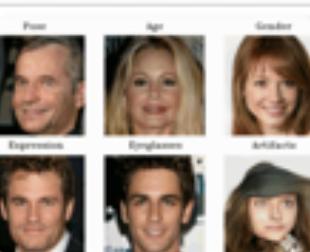


Image Processing Using Multi-Code GAN Prior

Jinjin Gu, Yujun Shen, Bolei Zhou

Computer Vision and Pattern Recognition (CVPR), 2020

[Paper] [Project Page] [Code]



Interpreting the Latent Space of GANs for Semantic Face Editing

Yujun Shen, Jinjin Gu, Xiaoou Tang, Bolei Zhou

Computer Vision and Pattern Recognition (CVPR), 2020

[Paper] [Project Page] [Code] [Demo Video]

Team



Principal Investigator

Bolei Zhou



Team Members

Yujun Shen



Ceyuan Yang



Jinjin Gu



Jiapeng Zhu

Conclusion

- Likelihood-based models: Autoregressive models
- Latent variable models: Variational Autoencoder
- Implicit generative models: Generative Adversarial Networks (GANs)
- Latent semantics in GANs

Conclusion

Unsupervised Learning

Generative
modeling

Self-supervised
feature learning